(DRAFT) Exploring Compositional Generalization (in ReCOGS_pos) by Transformers using Restricted Access Sequence Processing (RASP)

William Bruns

adde.animulis@gmail.com

Abstract

Humans rapidly generalize from a few observed examples and understand new combinations of words encountered if they are combinations of words recognized from different contexts, an ability called Compositional Generalization. Some observations contradict that Transformers learn systematic, compositional solutions to problems that generalize. The COGS benchmark (Kim and Linzen, 2020) reports zero percent accuracy for Transformer models on some structural generalization tasks. We use (Weiss et al., 2021)'s Restricted Access Sequence Processing (RASP), a Transformer-equivalent programming language, to prove by construction that a Transformer encoder-decoder can perform ReCOGS (Wu et al., 2024) systematically and compositionally while being flat and non-recursive/non-tree. Our RASP implementation suggests the reported "hardest split" obj-pp-to-subj-pp generalization in COGS may be a specific case of a general difficulty with ignoring prepositional-phrase-noun-phrases which are inserted in between two parts of speech with a relationship when the modified part-of-speech in the pair is on the left-side, predicting 100% of a certain category of error (Wu et al., 2024)'s baseline model makes on the split, and guides us to try testing the model on modifying recipient nouns on the right side of the verb in the v_dat_p2 form, on which the baseline Transformer performs similarly poorly on as predicted. The difficulty with the 'np v_dat_p2 np_det pp np np' modification and the previously reported subj pp modification difficulty are claimed by us to be incompatible with the Transformer learning a tree-based approach (which combines nonterminals as in 'np_det pp np -> np_pp -> np'), so we also check (Csordás et al., 2022)'s hypothesis that the number of Transformer layers should be at least the depth of the parse tree for a tree based solution and find no performance benefit to a couple of additional layers (beyond baseline of 2). Implementing our task in Restricted Access Sequence Processing immediately helped us

discover additional related failure modes of the baseline Encoder-Decoder Transformer, predict the details of the errors in the logical forms (what the mismatched index in the agent when the agent is on the left of the verb would be, 100% of the time) generated for the previously reported most difficult split, and may help us reason about why a model like (Wu et al., 2024) works with 2 layers for the ReCOGS task (compared with e.g. use of 24-layer BERT for NLP tasks in (Tenney et al., 2019)).

1 Introduction

Large pretrained language models based on the Transformers architecture compose seemingly fluent and novel text and are excellent few or zero shot learners (Brown et al., 2020), but some observations contradict that Transformers learn systematic, compositional solutions to problems that generalize, for example prepositional phrase generalizations in the ReCOGS (Wu et al., 2024) variant of the COGS task (Kim and Linzen, 2020).

Composition is important in learning. Consider a single nonterminal grammar expansion[1], '(noun phrase) (verb dative p2) (noun phrase) (noun phrase)', with three noun phrases all already expanded to np_det ("a" or "the" and "common noun") and a single verb. A possible substitution of terminals would be "a teacher gave the child a book", as would be "the teacher gave a child the book" (change of determiners), as would be "the manager sold a customer the car" (change of nouns and verb) and it would require $2^3 V_n^3 V_v$ examples where V_n is the vocab size for eligible common nouns and V_v is the vocab size for eligible verbs to see all the possible terminal substitutions . If the qualifying vocabulary is say of order of magnitude 100 words for the nouns and 10 for the verbs[2]

---

[1]COGS input sentences were actually generated by a probabilistic context-free grammar and this is a grammar expansion in their grammar. Words used in the example are within their vocabulary.

[2]In COGS the number of common nouns is over 400 and

that would come out order of magnitude 100 million examples. By contrast, if parts-of-speech and verb types are already known[3] it might take as few as one example to learn the new grammar pattern '(noun phrase) (verb dative p2) (noun phrase) (noun phrase)'. We will see in the results quantitatively how few training examples it actually takes to cover the grammar of problem we are solving in the sense of (Zeller et al., 2023).

For decades it was argued that connectionist models (i.e. neural networks) were somehow structurally incapable of compositional learning (Fodor and Pylyshyn, 1988). We know this must not be true given the success of the large language models, but seek to try a new toolkit to understand remaining failure modes.

We use (Weiss et al., 2021)'s Restricted Access Sequence Processing (RASP) language that can be compiled to concrete Transformer weights to prove by construction that a Transformer encoder-decoder can perform the ReCOGS (Wu et al., 2024) variant of the COGS (Kim and Linzen, 2020) task over the vocabulary and grammar of that task in a systematic, compositional way (length and recursion depth limited) as a rigorous starting point to investigating when Transformers might learn or not actually learn such compositional/systematic solutions. Our non-tree, non-recursive RASP implementation suggests the reported "hardest split" obj-pp-to-subj-pp generalization in COGS may be a specific case of a general difficulty with ignoring prepositional-phrase-noun-phrases which are inserted in between two parts of speech with a relationship when the modified part-of-speech in the pair is on the left-side, predicting successfully the details of the error (Wu et al., 2024)'s baseline model makes on the split, and guides us to try testing the model on a very different syntax predicted to have the same mechanism (and therefore the same difficulty), modifying recipient nouns on the right side of verb in the v_dat_p2 form, which it performs similarly poorly on as predicted by the RASP model. (However, trying to fix this by augmenting the data with these 'np v_dat_p2 np_det pp np np' examples does not yet show any crossover benefit

to both splits.) As the hypothesis that predicts the observed difficulty on 'np v_dat_p2 np_det pp np np' from the previously reported subj pp modification difficulty implies the Transformer is learning a flat, non-tree based approach, we also investigate (Csordás et al., 2022)'s hypotheses about the number of layers required to compositionally model a parsing problem in terms of the depth of the parse tree (compositional operations at each level suggests depth must exceed the height of the parse tree). We find no performance benefit to 3 or 4 layers instead of 2, further supporting that a tree based, recursive solution is NOT learned by the (Wu et al., 2024) baseline Encoder-Decoder Transformer for this simple grammar (consistent with our RASP model). This explains the challenge of generalizing on unseen prepositional phrase related modification related splits as arising from the baseline 2 to 4 layer[4] Encoder-Decoder Transformers not being able to leverage the grammar rule 'np_det pp np -> np_pp -> np' during learning and which requires them to instead actually observe more of the various prepositional phrase substitutions to learn them.

2 Prior Literature

(Wu et al., 2024) discuss the COGS benchmark which so far has stumped state of the art models with literally 0% accuracy on structural generalization splits, including length generalizations which seem analogous to (Clark et al., 2020) depth generalizations (perhaps surprising as for the simpler logical inferences problem in (Clark et al., 2020) they observed successful depth generalization), but (Wu et al., 2024) are able to get traction for the models in a modified form of the task they call ReCOGS by removing redundant symbols and using Semantic Exact Match instead of Exact Match to avoid forcing the model to predict arbitrary variable labeling in logical forms.[5]

(Zhou et al., 2023) apply (Weiss et al., 2021)'s

---

qualifying verbs in this case over 20

[3]that is if determiners ("a", "the") are understood to be equivalent, common nouns are already known ("teacher", "manager", "child", "customer", "book", "car") separately, qualifying verb dative verbs are already known ("gave", "sold"). Note (Tenney et al., 2019) report part-of-speech information is already tagged and in the very earliest layers of the 24-layer BERT large pre-trained language model.

---

[4]Deeper Transformers (as used in Large Language Models like GPT) may be a different story; (Tenney et al., 2019) analyze the extraction of increasingly complex information with layer depth in a 24-layer BERT.

[5]Note, the COGS/ReCOGS task is important because it is converting sentences to a representation of their meaning (in logical form where syntactically different but semantically identically sentences like "a boy painted the girl" or "the girl was painted by a boy" are written the same), and we want these models to be able to understand entirely the new sentences they will be encountering (probably immediately given the diversity of language) based on being familiar having observed pieces of them in different combinations in the past.

RASP language to explain some inconsistent findings regarding generalization and use RASP to predict exactly which cases of generalization come easily to Transformers and which do not. (Zhou et al., 2023) seem to reveal (Weiss et al., 2021) has provided the framework we seek by demonstrating how to apply RASP to Transformer decoders with intermediate steps, and even use it to learn how to modify difficult-to-learn tasks like Parity and long addition in seemingly incidental ways based on RASP analysis to make them readily learnable by Transformers in a compositional, length generalizing way!

Thus we apply techniques similar to (Zhou et al., 2023) and (Weiss et al., 2021) to ReCOGS which try to assess how (Wu et al., 2024)'s baseline Encoder-Decoder Transformers learn to extract semantics (logical form) from the syntax (grammatical form) of sentences and understand the prepositional phrase modification related generalization errors they are making.

3 Data

COGS (Kim and Linzen, 2020) and ReCOGS (Wu et al., 2024) datasets were used as provided by the repository associated with (Wu et al., 2024)[6]. The structural generalization splits (especially prepositional phrase Object-to-Subject modification) will be a focus.

The full grammar and vocabulary for COGS/ReCOGS English input sentences provided in the utilities associated with the IBM CPG project (Klinger et al., 2024)[7] were used in designing RASP solution and analyzing the ways in which this task could be learned. See Figure 0

4 Model

We used the RASP interpreter of (Weiss et al., 2021) to run our program. For RASP model design and details see Appendix 5. We use word-level tokens for all results in this paper.[8] Consistent with (Zhou et al., 2023) we use (Weiss et al., 2021)'s RASP originally used for modeling Transformer encoders to model an encoder-decoder in a causal way by feeding the autoregressive output back into

the program. We only have aggregations with non-causal masks when that aggregation (or without loss of generality just before the aggregation product is used to avoid multiplying everywhere) is masked by an input mask restricting it to the sequence corresponding to the input.[9]

For training Transformers from scratch with randomly initialized weights using gradient descent for comparison with RASP predictions, we use scripts derived from those provided by (Wu et al., 2024)[10].

See Appendix 8 - Model Detail.

5 Methods

We use the RASP (Weiss et al., 2021) intepreter[11] to evaluate our RASP programs[12].

Logical forms (LFs) which are scored by Semantic Exact Match[13] against ground truth.

For the RASP solution, due to time constraints, complement phrase examples e.g. "Emma noticed that ..." which it appears can be implemented in RASP, were excluded so the author could focus on full coverage of the remaining grammar within the two weeks[14].

We also measure grammar coverage of input examples supported by our RASP model against the full grammar of COGS/ReCOGS input sentences provided in the utilities of the IBM CPG project (Klinger et al., 2024)[15]

See Appendix 9 - Methods Detail.

---

[6] https://github.com/frankaging/ReCOGS

[7] https://github.com/IBM/cpg/blob/ c3626b4e03bfc681be2c2a5b23da0b48abe6f570 /src/model/cogs_data.py#L523

[8] We believe any solution at the word-level can be converted to a character-level token solution and that is not the focus of our investigation here (see Appendix 6 for proof of concept details on a character level solution not used here).

[9] An example the author has prepared of this is available at https://github.com/willy-b/learning-rasp/blob/ 16a8e154b025e91c8e56965a1d475e49f69ebdbd/recogs_examples_in_rasp.py .

[10] https://github.com/frankaging/ReCOGS/blob/ 1b6eca8ff4dca5fd2fb284a7d470998af5083beb/run_cogs.py and
https://github.com/frankaging/ReCOGS/blob/ 1b6eca8ff4dca5fd2fb284a7d470998af5083beb/model/encoder_decoder_hf.py

[11] provided at https://github.com/tech-srl/RASP/

[12] https://github.com/willy-b/learning-rasp/blob/ 16a8e154b025e91c8e56965a1d475e49f69ebdbd/word-level-pos-tokens-recogs-style-decoder-loop.rasp
with a demo at
https://github.com/willy-b/learning-rasp/blob/ 16a8e154b025e91c8e56965a1d475e49f69ebdbd/recogs_examples_in_rasp.py

[13] https://github.com/frankaging/ReCOGS/blob/ 1b6eca8ff4dca5fd2fb284a7d470998af5083beb/utils/train_utils.py and
https://github.com/frankaging/ReCOGS/blob/ 1b6eca8ff4dca5fd2fb284a7d470998af5083beb/utils/compgen.py

[14] Complement phrase support has been added in https://github.com/willy-b/learning-rasp/pull/7 , and the model scores 100% semantic exact match (95% confidence interval: 99.88% to 100.00%) on the full ReCOGS_pos Test set, and 100% exact string match (same confidence interval) on the dev set (test set results pending).

[15] https://github.com/IBM/cpg/blob/ c3626b4e03bfc681be2c2a5b23da0b48abe6f570/src/model/cogs_data.py#L523

6 Results

Restricted Access Sequence Processing - grammar coverage using a flat pattern matching approach (not tree-based and not recursive) and autoregressive decoder loop Given the COGS input sentences were generated as a probablistic context free grammar per (Kim and Linzen, 2020) using the full details put in Lark format by (Klinger et al., 2024) and converting it ourselves to a format compatible with (Zeller et al., 2023) (See Appendix 7 - Computing Grammar Coverage) , we use their TrackingGrammarCoverageFuzzer to generate the set of all expansions of the COGS grammar.

Before 75 examples of the ReCOGS training set, 100% grammar coverage is reached (lexical differences ignored and complement phrases omitted) (Zeller et al., 2023) (noting that if the model is not capable of learning certain expansions in the grammar such as 'np_det pp np -> np_pp -> np', it may need to see more variations to memorize individual cases instead):

% ReCOGS grammar coverage by # of examples in ReCOGS train set
(lexical differences ignored
as we show with Restricted Access Sequence Processing
the model can be built to treat all common nouns identically, all proper nouns identically, etc)
(complement phrases were excluded is why it does not reach 100%)

That shows if one already knows parts of speech and verb types for words one needs much less data.

Thus, we can be more efficient than using the ReCOGS training set for our RASP model built by hand (in real world pretraining or a GloVe embedding would be used to solve this step)[16] since we our solution uses a manual embedding via a dictionary mapping words to part-of-speech and verb-type, that ensures all words within a part of speech are treated identically. We generated 20 sentences which cover 100% of the grammar under those constraints (under the context free grammar, tree based assumption which turns out to be incorrect for prepositional phrases) per (Zeller et al., 2023):

```
"the girl was painted",
```

```
"a boy painted",
"a boy painted the girl",
"the girl was painted by a boy",
"a boy respected the girl",
"the girl was respected",
"the girl was respected by a boy",
"the boy grew the flower",
"the flower was grown",
"the flower was grown by a boy",
"the scientist wanted to read",
"the guest smiled",
"the flower grew",
"ella sold a car to the customer",
"ella sold a customer a car",
"the customer was sold a car",
"the customer was sold a car by ella",
"the car was sold to the customer by ella",
"the car was sold to the customer",

"a boy painted the girl in a house"
```

The first 19 of those sentences are present in our RASP program code[17] and each correspond to a set of RASP operations corresponding to attention operations in a Transformer to match a template corresponding to that sentence type[18]. Those 19 examples reflect the only rules for handling non-prepositional grammar rules[19] present in our RASP solution (which gets approximately 100% on the (Wu et al., 2024) ReCOGS_pos test set, see next results section).

Restricted Access Sequence Processing - semantic exact match

The *Restricted Access Sequence Processing program scored 100% Semantic Exact Match (no missed examples) (95% confidence interval by Beta distribution of 99.9% to 100%, n=2720) on the ReCOGS_pos test set*[20] (excluding complement phrase examples).[21].

---

[16](Tenney et al., 2019) confirm BERT, a pre-trained Transformer model in wide use, has part-of-speech information available at the earliest layers.

[17] https://github.com/willy-b/learning-rasp/blob/dca0bc6689b0454b75e5a46e77ffe66566ca7661/word-level-pos-tokens-recogs-style-decoder-loop.rasp#L568

[18]see also "Appendix 5 - Restricted Access Sequence Processing word-level (post-embedding) token program/model design"

[19]To handle prepositional phrases in a flat solution, we find it necessary even on the training data to add a rule that ignores noun phrases preceded by a prepositonal phrase (ignore "pp np") when searching for noun indexes to report in relationships (agent, theme, recipient, etc), and we loosen verb type templates to allow a gap for any prepositional phrase to be inserted (this was not done in every possible case manually within 2 weeks which is why we do not get 100% in that one split, see next results section). We shall see encountering this issue in RASP and the grammar analysis suggesting a non-tree solution leads us to be able to predict 100% of a certain category of errors a baseline Wu et al 2023 Encoder-Decoder Transformer makes in an upcoming results section.

[20]https://raw.githubusercontent.com/frankaging/ReCOGS/1b6eca8ff4dca5fd2fb284a7d470998af5083beb/recogs_positional_index/test.tsv ,
see https://colab.research.google.com/drive/1vsuY6YW3F-DigpH2bCdU1b5vRPouyig-

[21]Complement phrase support has been added since this paper was drafted in https://github.com/willy-b/learning-rasp/pull/7 , and the model scores 100% semantic exact match (95% confidence interval: 99.88% to 100.00%) on the full

The RASP model scored 99.3% Semantic Exact Match accuracy (95% confidence interval: 98.7% to 99.6%)[22] on the first 1,521 examples of the ReCOGS_pos generalization set[23] (excluding complement phrase examples and with prepositional phrase recursion split analyzed separately below).

See "Appendix 0 - Full RASP Program ReCOGS gen set Semantic Exact Match results by split" for all splits.

Restricted Access Sequence Processing - prepositional phrase recursion (tail recursive) with a non-tree, non-recursive approach using the decoder loop[24]

RASP model ReCOGS pp_recursion gen split scores on first 83 examples were 100% (95% confidence interval: 95.7% to 100.0%; no missed examples have yet been observed in n=83)[25]

Wu et al 2023 Encoder-Decoder Transformer from scratch baselines (ReCOGS_pos)

Wu et al 2023's baseline Encoder-Decoder Transformer on ReCOGS_pos had an overall score of 88.55% +/- 1.87% Semantic Exact Match accuracy (sample +/- std, n=20) with a 95% confidence

---

ReCOGS_pos Test set, and 100% exact string match (same confidence interval) on the dev set (test set results pending), see PR for details

[22]*The only missed examples were in the obj_pp_to_subj_pp split, which scored 89.9% (95% confidence interval: 82.7% to 94.9%)*, due to two week time constraint and limitations on the author's ability to input all the grammar defined prepositional phrase rules in within that time (should not be taken as an indicator this is an upper bound for a similar RASP program, I expect 100% is feasible like the other splits).

[23] https://raw.githubusercontent.com/frankaging/ReCOGS/1b6eca8ff4dca5fd2fb284a7d470998af5083beb/recogs_positional_index/gen.tsv , see https://colab.research.google.com/drive/1Uw3gn9PXpAs6rOurL5_fLvo4o6_ewMeT .
The RASP interpreter is not that fast and our program is not optimized, so we could not evaluate on the entire 21K example gen set within this short project.
Thus we report on the first 1,521 examples here.

[24]The grammar includes two (tail) recursive aspects, prepositional phrase recursion, covered here, and complement phrase recursion (I argue can be handled identically to prepositional phrase recursion so just covered one within this 2 week project).
The prepositional phrase recursion comes from the following COGS input grammar rules: 'np -> np_det | np_prop | np_pp' and 'np_pp -> np_det pp np'.
Thus np can be expanded in an unbounded way as follows: 'np -> (np_det pp np) -> np_det pp (np_det pp np) -> np_det pp np_det pp (np_det pp np)' and so on.
However, one sees this is tail recursion and can be handled by a loop that just appends 'np_det pp' until the final 'np' is not 'np_pp'.

[25] Notebook: https://colab.research.google.com/drive/19mpX9kNbJoH-70rsbaXK2pWD19CbexDW

---

interval for the sample mean when n=20 of 87.73% to 89.37%.

Their Semantic Exact Match score on the extremely difficult obj_pp_to_subj_pp split was 19.7% +/- 6.1% Semantic Exact Match accuracy (sample +/- std, n=20) with 95% confidence interval for the sample mean with n=20 of 17.0% to 22.4% .

Wu et al 2023 Encoder-Decoder baseline 2-layer Transformer does not improve when adding 1 or 2 additional layers[26](even allowing parameter count to increase)[27]

3-layer Wu et al 2023 Encoder-Decoder: 16.1% +/- 1.7% Semantic Exact Match (sample mean +/- std) with 95% confidence interval for sample mean (n=8) of 14.9% to 17.3%.

4-layer Wu et al 2023 Encoder-Decoder: 19.2% +/- 4.4% Semantic Exact Match (sample mean +/- std) with 95% confidence interval for sample mean (n=5) of 15.4% to 23.1%.

Error Analysis for Wu et al 2023 baseline Encoder-Decoder Transformer on obj_pp_to_subj_pp split

To assess whether the mechanism of the errors in the obj_pp_to_subj_pp generalization split was what we expect from the Restricted Access Sequence Processing program, and were due to the Transformer naively matching the closest noun (now the prepositional phrase noun) instead of the original subject noun when the noun to the left of the verb was modified by a prepositional phrase, adding a prepositional phrase noun closer than the target noun, we analyzed the specific errors the Wu et al 2023's Transformer made.

The errors from a run of the baseline Wu et al 2023 Encoder-Decoder Transformer trained on their ReCOGS_pos train.tsv and tested on their unmodified gen.tsv were analyzed. All the input sentences and output logical forms as well as the ground truth logical forms were logged during the run. The input sentences were parsed by the Lark parser[28] against the COGS input grammar which allowed categorizing each sentence by its verb type

---

[26]This is consistent with the flat, non-tree solution we argue for in ths paper, that e.g. cannot learn to combine 'np_det pp np -> np_pp -> np', makes the predicted mistakes in the subj pp when the agent is left of the verb, and does poorly on our new v_dat_p2_pp_moved_to_recipient split (see discussion).

[27]Since no improvement was observed, we did not run the costly experiments to increase the layers while controlling the parameter count (which would be a follow up to distinguish if the improvement was from the layer increase or the parameter increase).

[28]https://github.com/lark-parser/lark

[29]. The author manually inspected each of verb type patterns and categorized them by the position of the agent and theme relative to the verb (see "Appendix 3 - Error Analysis - parsing sentences with Lark and tagging sentences as agent left-of-verb or not") as used below.

For the difficult obj_pp_to_subj_pp split, considering only single verb cases (except v_inf_taking to v_inf which was included) and ignoring sentences with complement phrases, of the sentences with a logical form failing Semantic Exact Match, 30.4% of the errors were single relationship errors in one of the agent, theme, or recipient relationships.[30]

For the 30.4% of errors in this split which were single relationship errors, 75.4% of these single relationship errors had the agent on the left side. *Critically, 100% (95% confidence interval by Beta distribution: 98.6% to 100%) of the single relationship errors with the agent on the left side were in the agent part of the logical form and 100% (95% confidence interval: 98.6% to 100%) of those single relationship errors in the agent (258 examples) were that the agent index was accidentally assigned to the modifying prepositional phrase noun instead of the original agent noun (as predicted by the RASP analysis).*

Wu et al 2023 Encoder-Decoder Transformer on new v_dat_p2 pp moved to recipient (from theme) split - as hard as hardest previous generalization split

As the Restricted Access Sequence Processing program predicted the 'np v_dat_p2 np pp np np'[31] prepositional phrase modification (which involves the recipient instead of the subject so is a distinct check of our hypothesis) would require learning to ignore the distractor "pp np" same as required for the obj_pp_to_subj_pp split, we predicted that a new split we introduce "v_dat_p2_pp_moved_to_recipient" would also be difficult for the Transformer. To test this, Wu et al

---

[29]Code to analyze the errors is at: https://colab.research.google.com/drive/1MiiEAchmaGulsTwNHs98-ill-UM7TK3i

[30]Of the single relationship errors, we categorized them by a description of the position of both the agent and theme relative to the verb in that sentence (agent was considered to be either left OR "right or middle"; theme could be left, right, or middle) and what relationship had the error. Based on the motivation discussed in the analysis section, we focused on the case where the agent was on the left and the theme was in any position.

[31]Strictly speaking we only do 'np v_dat_p2 np_det pp np np' as per the grammar 'np_prop' cannot precede a prepositional phrase

2023's baseline Encoder-Decoder Transformer was trained with default data (ReCOGS_pos train.tsv) and tested on modified v_dat_p2 pp training examples where only the word order was changed to move the prepositional phrase from the theme to the recipient (logical form properly updated see Appendix 4 for all examples). *The baseline Wu et al 2023 Encoder-Decoder Transformer was only able to achieve a Semantic Exact Match (sample mean +/- sample std) of 13% +/- 15.6% (n=10 samples) with a 95% confidence interval for the sample mean when n=10 of 4% to 23%.* Thus, this new split we introduce here as v_dat_p2_pp_moved_to_recipient is as difficult or perhaps more difficult than the previous reported "hardest split" obj_pp_to_subj_pp.

Wu et al 2023 Encoder-Decoder Transformer trained with data augmented with v_dat_p2 pp moved to recipient (from theme) does NOT improve obj_pp_to_subj_pp performance

Wu et al 2023's baseline Encoder-Decoder Transformer was trained with default data (ReCOGS_pos train.tsv) but with additionally the same modified v_dat_p2 pp training examples used for the "v_dat_p2_pp_moved_to_recipient" split (non-subject recipient modified with prepositional phrase, so nonoverlapping with subj_pp) above on which it performed poorly, then tested on the standard prepositional modification generalization split "obj_pp_to_subj_pp", after which it achieved 22% +/- 6.7% Semantic Exact Match (sample mean +/- std, n=10) with 95% confidence interval for sample mean n=10 of 17.9% to 26.1% (not significantly different than Wu et al 2023's baseline by one-tailed Welch's unequal variances t-test).

7 Analysis

Our RASP model of a Transformer Encoder Decoder, without tree-based or recursive aspects scored 100% in semantic exact match accuracy on the (Wu et al., 2024) test set (n=2720), and on the generalization data scored 100% in all but one category (see above) without explicit rules in the RASP program to handle them. This includes 100% accuracy on the prepositional phrase recursion generalization split up to depth 12 (n=83 examples), without any hardcoded prepositional phrase expansion shortcuts added[32]. The RASP program only made a

---

[32]a single rule applies to all depths; the only limit on length generalization is the RASP interpreter and a simple to extend positional encoding which only handles sentences up to 50

significant number of errors on obj_pp_to_subj_pp which scored only 89.9% (95% confidence interval: 82.7% to 94.9%) Semantic Exact Match accuracy[33] much better than (Wu et al., 2024) baseline Encoder-Decoder Transformers which only scored 19.7% +/- 6.1% with 95% confidence interval for the sample mean with n=20 of 17.0% to 22.4% .

Note that complement phrase recursion was considered redundant with prepositional phrase recursion as they are both tail recursion cases that can be handled by the decoder loop and omitted to fit the project within 2 weeks (will be done as follow up).

Thus, we demonstrated by construction using the Restricted Access Sequence Processing language which can be compiled to concrete Transformer weights that theoretically a Transformer Encoder-Decoder can solve the COGS input to ReCOGS_pos logical form translation in a systematic, compositional, and length generalizing way.

To assess whether the mechanism of the errors in the obj_pp_to_subj_pp generalization split was what we expect from the Restricted Access Sequence Processing program, and were due to the Transformer naively matching the closest noun[34] (now the prepositional phrase noun) instead of the original subject noun when the noun to the left of the verb was modified by a prepositional phrase, adding a prepositional phrase noun closer than the target noun, we analyzed the specific errors the Wu et al 2023's Transformer made. Thus we analyzed the case where the noun being modified to the left corresponded to the agent (the most common case, with active tense verbs in the subject pp modification split). In 100% of the 258 cases where there was a single relationship error and the agent on the left the error was specifically in the agent relationship and the mistake was as predicted: the model mistook the inserted prepositional phrase noun to be the agent when it was inserted on the left of the verb closer than the actual agent noun (occurs when agent being modified by prepositional phrase is on left due to asymmetry of pp modifications adding on the right side expanding to the right).

---

words due to a map that only covers numbers up to 50 but can be easily expanded by literally adding dummy entries like "'51':0,'52':0" and so on to a map in one place

[33]and only less than 100% on this split because within the 2 week project time the author could not cover all the prepositional grammar expansion rules in coding the rule-based RASP solution by hand

[34]See Figure 1 at the end.

Our explanation could have been refuted by other single relationship errors occurring as frequently as the agent, indicating general model confusion (mixing up agent and theme, not just agent relationships) and/or when making an agent error, the model could have simply put nonsense indices or referred to any other position other than the closest noun position displaced by the pp noun to refute our hypothesis.

The simple mechanism of closest noun being displaced when modifying a noun to the left with a prepositional phrase can also be checked by making a prediction on a completely different syntax affected by the same issue: the 'np v_dat_p2 np pp np np'[35] prepositional phrase modification (which involves the recipient relationship being modified instead of the subject and/or agent so is a distinct check of our hypothesis). According to our Restricted Access Sequence Processing program as well the 'np v_dat_p2 np pp np np' generalization requires learning the exact same rule as the subj pp generalization: to ignore the distractor "pp np". Thus we predicted that a new split we introduce, "v_dat_p2_pp_moved_to_recipient", would also be difficult for the baseline Wu et al 2023 Encoder-Decoder Transformer. To test this, Wu et al 2023's baseline Encoder-Decoder Transformer was trained with default data (ReCOGS_pos train.tsv) and tested on modified v_dat_p2 pp training examples where only the word order was changed to move the prepositional phrase from the theme to the recipient (logical form properly updated see Appendix 4 for all examples) and we indeed found that this was as hard or harder than the previous most difficult split analyzed above, the 'obj_pp_to_subj_pp' split.

Note that if the Encoder-Decoder Transformer were to learn a tree-based or recursive representation, it would also be predicted that the "v_dat_p2_pp_moved_to_recipient" would not be any harder than when the pp modification is on the theme, as 'np v_dat_p2 np_det pp np np' can be transformed by the recursive grammar rule 'np_det pp np -> np_pp -> np' to 'np v_dat_p2 np np' on which it is already trained and has good performance, whereas we observe "v_dat_p2_pp_moved_to_recipient" is instead as hard as the hardest previously reported generalization split.

We know that Transformers only move informa-

---

[35]Again, precisely we only do 'np v_dat_p2 np_det pp np np' as per the grammar 'np_prop' cannot precede a prepositional phrase

tion between positions in a sequence at each layer boundary via cross/self-attention. (Csordás et al., 2022) argues the Transformer layers should be as deep as the deepest data dependency in the computatonal graph of the problem, in our case the parse tree[36]. Maybe then, (Wu et al., 2024) baseline Encoder-Decoder is not constrainted to learn a flat, non-tree model with these characteristic errors and with more layers it would learn to recursively combine 'np_det pp np -> np_pp -> np' (to some fixed depth at least, probably limited by the depth) and perform better on prepositional phrase related splits[37].

However, training a (Wu et al., 2024) baseline Encoder-Decoder Transformer from scratch we found no benefit to 3 or 4 layers instead of 2.

Taken together, these results and the grammar coverage analysis suggest we may interpret the poor performance on generalizing on unseen prepositional phrase related modification related splits as arising from the baseline 2 to 4 layer Encoder-Decoder Transformers learning a flat representation (non-tree, non-recursive) that is not being able to leverage the grammar rule 'np_det pp np -> np_pp -> np' during learning and which requires them to instead actually observe more of the various prepositional phrase substitutions to learn them.

It is hoped this candidate explanation for the problem, will in future work assist in fixing it with clever training data augmentations[38], or curriculum learning, or architectural changes.

---

[36]"the network should be sufficiently deep, at least as deep as the deepest data dependency in the computational graph built from elementary operations (e.g., in the case of a parse tree, this is the depth of the tree)"

[37]this is not a very scalable approach as we must make the network deeper to handle longer prepositional chains instead of just looping

[38]We note we attempted in the course of this work one simple data augmentation (see Methods) which did not improve performance on the obj-pp-to-subj-pp split which was to attempt to teach the network to ignore the distracting "pp np" (a rule we had to add to the RASP, see Appendix 5 and Model sections) which along with templates with dynamic size gaps (match any occurrence of something in the future rather than a fixed length away) allows our RASP model to achieve high obj-pp-to-subj-pp scores despite being a non-tree solution. A simple explanation for the augmentation not fixing the problem is that our augmentation would not help the Transformer learn to alter the templates, only to ignore "pp np" when outputting nouns in relationships, even in the error analysis where we predicted the location and value of 100% of agent-left-side-in-input single point errors, only 30% of examples had errors in one position, and the rest may have to do with not detecting the right verb type template, not just matching the agent, recipient, theme nouns.

8   Conclusion

Implementing our task in Restricted Access Sequence Processing immediately helped us discover additional related failure modes (e.g. new "v_dat_p2_pp_moved_to_recipient" split) of the baseline Encoder-Decoder Transformer, predict the details of the errors in the logical forms and may help us reason about why a model like (Wu et al., 2024) can work with 2 layers for the ReCOGS task (can be solved by a flat, non-tree approach), and recommend others to consider to use RASP to understand Transformer behavior even for more complicated tasks like ReCOGS.

Known Project Limitations

The Restricted Access Sequence Processing code is not optimized. Cannot yet predict attention heads and layers required from the select and aggregate operations performed like the RASP authors (Weiss et al., 2021) were able to do with their problems.

Our RASP evaluation is slow.

Grammar coverage (Zeller et al., 2023) is only valid when the expansions are rules your model can learn.[39] We specifically made use of this limitation in this paper but still caution anyone about it who might just take the grammar coverage metric away by itself.

The error analysis of the (Wu et al., 2024) baseline Encoder-Decoder Trnsformer on the obj_pp_to_subj_pp split (predicting the 258 point errors made when agent was on the left side of the verb) relies on a single model's behavior though strongly statistically significant, and should be replicated.

Much deeper Transformer networks may be learning a tree-based grammar representation and not suffer from the predicted generalization issues observed in (Wu et al., 2024)'s baseline 2-layer Transformer and our intentionally non-tree RASP model.[40]

---

[39]If for example, as with our flat RASP model by design or as we hypothesize for (Wu et al., 2024)'s baseline Encoder-Decoder Transformer, the model cannot or will not learn the rule 'np_det pp np -> np_pp -> np' which recursively replaces noun phrases modified by prepositional phrases with a noun phrase, then grammar coverage will assume any prepositional phrase exposure is sufficient, which is evidently not true given the errors on prepositional phrase modification generalization splits reported here and by (Wu et al., 2024), (Kim and Linzen, 2020).

[40]Nothing explored here rules that out and there is plenty of evidence outside the COGS task-related literature suggesting this will be the case (e.g. (Tenney et al., 2019) whose 24-layer BERT model which seems to handle "POS tagging, parsing,

References

Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language models are few-shot learners.

Peter Clark, Oyvind Tafjord, and Kyle Richardson. 2020. Transformers as soft reasoners over language.

Róbert Csordás, Kazuki Irie, and Jürgen Schmidhuber. 2022. The neural data router: Adaptive control flow in transformers improves systematic generalization.

Jerry A. Fodor and Zenon W. Pylyshyn. 1988. Connectionism and cognitive architecture: A critical analysis. *Cognition*, 28(1):3–71.

Najoung Kim and Tal Linzen. 2020. COGS: A compositional generalization challenge based on semantic interpretation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 9087–9105, Online. Association for Computational Linguistics.

Tim Klinger, Luke Liu, Soham Dan, Maxwell Crouse, Parikshit Ram, and Alexander Gray. 2024. Compositional program generation for few-shot systematic generalization.

Ian Tenney, Dipanjan Das, and Ellie Pavlick. 2019. Bert rediscovers the classical nlp pipeline.

Gail Weiss, Yoav Goldberg, and Eran Yahav. 2021. Thinking like transformers.

Zhengxuan Wu, Christopher D. Manning, and Christopher Potts. 2024. Recogs: How incidental details of a logical form overshadow an evaluation of semantic interpretation.

Andreas Zeller, Rahul Gopinath, Marcel Böhme, Gordon Fraser, and Christian Holler. 2023. Grammar coverage. In *The Fuzzing Book*. CISPA Helmholtz Center for Information Security. Retrieved 2023-11-11 18:18:06+01:00.

Hattie Zhou, Arwen Bradley, Etai Littwin, Noam Razin, Omid Saremi, Josh Susskind, Samy Bengio, and Preetum Nakkiran. 2023. What algorithms can transformers learn? a study in length generalization.

NER, semantic roles, then coreference") and it is a goal that in future work to try to find data augmentation tricks or curriculum learning approaches that can force the model into such a learning mode for the COGS task at the minimum number of layers.

9   Notes

No AI tools were used by the author in the preparation of this manuscript with the exception of anything used in the backend by Google Scholar searches for literature and citations and Google searches for related material. AI writing aids were not used.

10 Figure 0: Translating sentences to ReCOGS logical form, syntax vs semantics
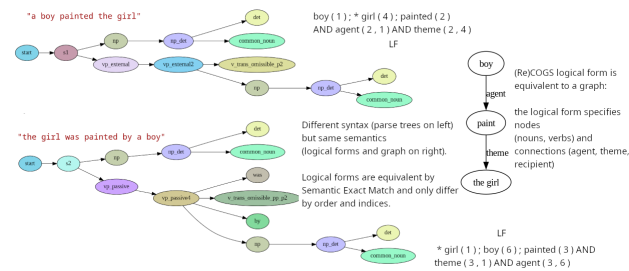
11 Figure 1: Explaining Subj PP generalization difficulties in ReCOGS for Transformers as non-tree model struggling with PP NP distractors when modifying NPs with relationships to the right
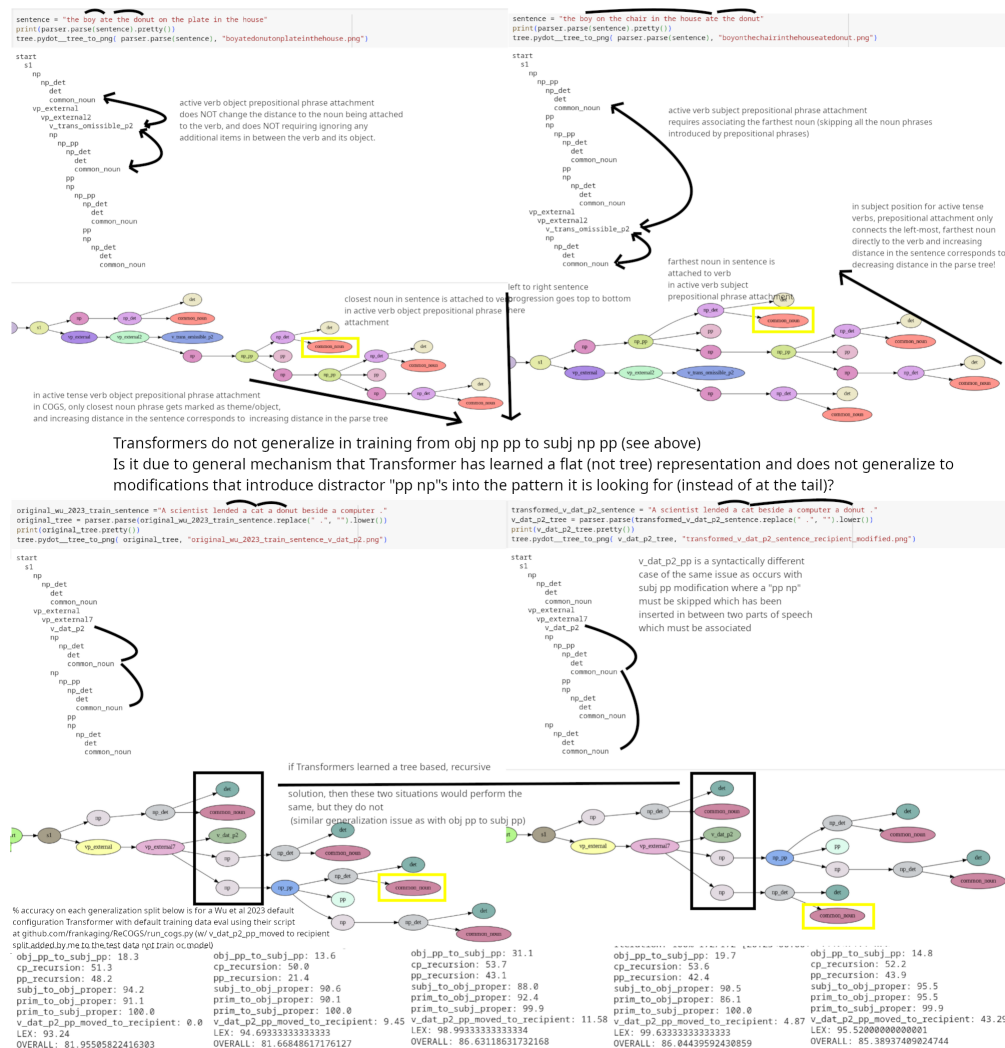
12 Appendix 0 - RASP Program ReCOGS gen set Semantic Exact Match results by split

Only 1521 rows are shown as the RASP program does not run quickly and the entire gen set was not able to be evaluated yet.

```
99.27679158448389% or 1510.0 out of 1521 (95% confidence interval: 98.7\% to 99.6\%)

Semantic Exact Match % by category:                                          Semantic Exact Match
Category
active_to_passive                                    100.000000
do_dative_to_pp_dative                               100.000000
obj_omitted_transitive_to_transitive                 100.000000
obj_pp_to_subj_pp                                     89.908257 (95% confidence interval: 82.7% to 94.9%)
obj_to_subj_common                                   100.000000
obj_to_subj_proper                                   100.000000
only_seen_as_transitive_subj_as_unacc_subj           100.000000
only_seen_as_unacc_subj_as_obj_omitted_transiti...   100.000000
only_seen_as_unacc_subj_as_unerg_subj                100.000000
passive_to_active                                    100.000000
pp_dative_to_do_dative                               100.000000
prim_to_inf_arg                                      100.000000
prim_to_obj_common                                   100.000000
prim_to_obj_proper                                   100.000000
prim_to_subj_common                                  100.000000
prim_to_subj_proper                                  100.000000
subj_to_obj_common                                   100.000000
subj_to_obj_proper                                   100.000000
unacc_to_transitive                                  100.000000

Counts by category:
Category
active_to_passive                                     77
do_dative_to_pp_dative                                88
obj_omitted_transitive_to_transitive                 100
obj_pp_to_subj_pp                                     109
obj_to_subj_common                                    113
obj_to_subj_proper                                    100
only_seen_as_transitive_subj_as_unacc_subj             5
only_seen_as_unacc_subj_as_obj_omitted_transiti...     5
only_seen_as_unacc_subj_as_unerg_subj                  1
passive_to_active                                     95
pp_dative_to_do_dative                               107
prim_to_inf_arg                                       81
prim_to_obj_common                                   109
prim_to_obj_proper                                    83
prim_to_subj_common                                   86
prim_to_subj_proper                                   98
subj_to_obj_common                                    74
subj_to_obj_proper                                    94
unacc_to_transitive                                   96
```

13 Appendix 1 - Vocabulary and Grammar

```
start:
    s1 | s2 | s3 | s4 | vp_internal
s1: np vp_external
s2: np vp_passive
s3: np vp_passive_dat
s4: np vp_external4
vp_external:
    v_unerg | v_trans_omissible_p1
    | vp_external1 | vp_external2
    | vp_external3 | vp_external5
    | vp_external6 | vp_external7
vp_external1: v_unacc_p1 np
vp_external2:
    v_trans_omissible_p2 np
vp_external3:
    v_trans_not_omissible np
vp_external4:
    v_inf_taking to v_inf
vp_external5:
    v_cp_taking that start
vp_external6:
    v_dat_p1 np pp_iobj
vp_external7:
    v_dat_p2 np np
vp_internal: np v_unacc_p2
vp_passive: vp_passive1 | vp_passive2
            | vp_passive3 | vp_passive4
            | vp_passive5 | vp_passive6
            | vp_passive7 | vp_passive8
vp_passive1:
    was v_trans_not_omissible_pp_p1
vp_passive2:
    was v_trans_not_omissible_pp_p2
          by np
vp_passive3:
    was v_trans_omissible_pp_p1
vp_passive4:
    was v_trans_omissible_pp_p2 by np
vp_passive5:
    was v_unacc_pp_p1
vp_passive6:
    was v_unacc_pp_p2 by np
vp_passive7:
    was v_dat_pp_p1 pp_iobj
vp_passive8:
    was v_dat_pp_p2 pp_iobj by np
>vp_passive_dat:
    vp_passive_dat1
            | vp_passive_dat2
vp_passive_dat1:
    was v_dat_pp_p3 np
vp_passive_dat2:
    was v_dat_pp_p4 np by np
np:
    np_prop | np_det | np_pp
np_prop: proper_noun
np_det: det common_noun
np_pp: np_det pp np
pp_iobj: to np
det: "the" | "a"
pp: "on" | "in" | "beside"
was: "was"
by: "by"
to: "to"
that: "that"
common_noun:
    "girl" | "boy"
    | "cat" | "dog" | ...
proper_noun:
    "emma" | "liam"
    | "olivia" | "noah"
    | ...
v_trans_omissible_p1:
    "ate" | "painted" | "drew"
    | "cleaned" | ...
v_trans_omissible_p2:
    "ate" | "painted" | "drew"
    | "cleaned"  | ...
v_trans_omissible_pp_p1:
    "eaten" | "painted" | "drawn"
    | "cleaned" | ...
v_trans_omissible_pp_p2:
    "eaten" | "painted" | "drawn"
    | "cleaned" | ...
v_trans_not_omissible:
    "liked" | "helped" | "found"
    | "loved" | ...
v_trans_not_omissible_pp_p1:
    "liked" | "helped" | "found"
    | "loved"  | ...
v_trans_not_omissible_pp_p2:
    "liked" | "helped" | "found"
    | "loved" | ...
v_cp_taking:
    "liked" | "hoped" | "said"
```

```
    | "noticed" | ...
v_inf_taking:
    "wanted" | "preferred" | "needed"
    | "intended" | ...
v_unacc_p1:
    "rolled" | "froze" | "burned"
    | "shortened" | ...
v_unacc_p2:
    "rolled" | "froze" | "burned"
    | "shortened" | ...
v_unacc_pp_p1:
    "rolled" | "frozen" | "burned"
    | "shortened" | ...
v_unacc_pp_p2:
    "rolled" | "frozen" | "burned"
    | "shortened" | ...
v_unerg:
    "slept" | "smiled" | "laughed"
    | "sneezed" | ...
v_inf:
    "walk" | "run" | "sleep"
    | "sneeze" | ...
v_dat_p1:
    "gave" | "lended" | "sold"
    | "offered" | ...
v_dat_p2:
    "gave" | "lended" | "sold"
    | "offered" | ...
v_dat_pp_p1:
    "given" | "lended" | "sold"
    | "offered" | ...
v_dat_pp_p2:
    "given" | "lended" | "sold"
    | "offered" | ...
v_dat_pp_p3:
    "given" | "lended" | "sold"
    | "offered" | ...
v_dat_pp_p4:
    "given" | "lended" | "sold"
    | "offered" | ...
```

14 Appendix 2 - RASP for relation right index ignoring distractor "pp np"

```
pp_sequence = \
indicator(pos_tokens == 2);
pp_one_after_mask = \
select(pp_sequence, 1, ==) and \
select(indices+1, indices, ==);

pp_one_after_sequence = \
aggregate(pp_one_after_mask, 1);
pp_one_after_mask = \
select(pp_one_after_sequence, 1, ==) and \
select(indices, indices, ==);

pp_two_after_mask = \
select(pp_sequence, 1, ==) and \
select(indices+2, indices, ==);

pp_two_after_sequence = \
aggregate(pp_two_after_mask, 1);
pp_two_after_mask = \
select(pp_two_after_sequence, 1, ==) and \
select(indices, indices, ==);

np_det_diag_mask = \
select(aggregate(np_det_mask, 1), 1, ==) and \
select(indices, indices, ==);

np_prop_diag_mask = \
select(aggregate(np_prop_mask, 1), 1, ==) and \
select(indices, indices, ==);

no_pp_np_mask = \
1 - aggregate((pp_one_after_mask and np_prop_diag_mask) or \
(pp_two_after_mask and np_det_diag_mask), 1);

nps_without_pp_prefix_indices = \
selector_width(select(NOUN_MASK*no_pp_np_mask, 1, ==) and \
select(indices, indices, <=))*NOUN_MASK*no_pp_np_mask;

left_idx = \
aggregate(select(indices, left_idx_in_nps_zero_based, ==), input_indices_sorted);
right_idx = \
aggregate(select(nps_without_pp_prefix_indices, after_intro_idx, ==), indices); # <--
```

15 Appendix 3 - Error Analysis - parsing sentences with Lark and tagging sentences as agent left-of-verb or not

```
# full error analysis code at https://colab.research.google.com/drive/1MiiEAchmaGulsTwNHs98-ill-UM7TK3i

# used grammar string defined in Appendix 1.
parser = Lark(grammar, start='start')

# 1st NP agent verbs (non CP)
# "v_trans_omissible_p1": "agent",
# "v_trans_omissible_p2": "agent",
# "v_trans_not_omissible": "agent",
# "v_cp_taking": "agent",
# "v_inf_taking": "agent",
# "v_unacc_p1": "agent",
# "v_unerg": "agent",
# "v_inf": "agent",
# "v_dat_p1": "agent",
# "v_dat_p2": "agent",
agent_left_of_verb_verb_type_set = set(["v_trans_omissible_p1", "v_trans_omissible_p2",
"v_trans_not_omissible", "v_cp_taking", "v_inf_taking",
"v_unacc_p1", "v_unerg", "v_inf", "v_dat_p1", "v_dat_p2"])

theme_left_of_verb_verb_type_set = set(
    ["v_trans_omissible_pp_p1",
     "np v_unacc_p2",
     "v_unacc_pp_p1",
     "v_unacc_pp_p2",
     "v_trans_omissible_pp_p2",
     "v_trans_not_omissible_pp_p1",
     "v_trans_not_omissible_pp_p2",
     "v_dat_pp_p1",
     "v_dat_pp_p2"
     ]) # fill this out

theme_right_of_verb_verb_type_set = set([
    "v_unacc_p1",
    "v_trans_omissible_p2",
    "v_trans_not_omissible",
]) # fill this out
theme_middle_of_dative_verb_type_set = set(["v_dat_pp_p4", "v_dat_p1"])

def get_verbs(lark_tree_root):
  nodes = [lark_tree_root]
  verbs = []
  while len(nodes) > 0:
    node = nodes[-1]
    nodes = nodes[:-1]
    node_type = node.data[:]
    if node_type[:2] == 'v_':
      verbs.append(node_type)
    for child in node.children:
      # it is a tree, no need to check for revisits
      nodes.append(child)
  return verbs

def get_theme_side(lark_tree_root):
    verb_type = get_verbs(lark_tree_root)[0]
    if verb_type in theme_right_of_verb_verb_type_set:
        return "right"
    elif verb_type in theme_left_of_verb_verb_type_set:
        return "left"
    elif verb_type in theme_middle_of_dative_verb_type_set:
        return "middle"
    return None

def get_agent_side(lark_tree_root):
    verb_type = get_verbs(lark_tree_root)[0]
    if verb_type != None and verb_type not in agent_left_of_verb_verb_type_set:
        return "right or middle"
    elif verb_type in agent_left_of_verb_verb_type_set:
        return "left"
    return None
```

16 Appendix 4 - v_dat_p2 recipient pp-modification for generalization assessment and data augmentation attempt

We test generalization by the (Wu et al., 2024)'s default Transformer which has been trained on 'np v_dat_p2 np np pp np' but not 'np v_dat_p2 np pp np np' prepositional modifications. The following 328 examples were derived[41] from the existing

https://github.com/frankaging/ReCOGS/blob/1b6eca8ff4dca5fd2fb284a7d470998af5083beb/recogs_positional_index/train.tsv ,

by modifying 328 existing single-pp v_dat_p2 lines in train.tsv to simply move the prepositional phrase from the 3rd NP (theme) in the 'np v_dat_p2 np np' (agent, recipient, theme) to the 2nd NP (recipient), e.g. copying and modifying the line "Liam gave the monkey a chalk in the container ." to "Liam gave the monkey in the container a chalk .".

So all the words and the grammar are otherwise familiar. This is similar to the existing 'obj_pp_to_subj_pp' generalization (Wu et al., 2024) reports on. All modified rows available in the notebook link in the footnote.

---

[41]Notebook: https://colab.research.google.com/drive/1IDs0EwIMp2wtLHk4KqnuGhuT3G14QEG1?usp=sharing

17 Appendix 5 - Restricted Access Sequence Processing word-level (post-embedding) token program/model design

You can run a demo and see the autoregressive output

(or just visit https://colab.research.google.com/drive/1FS4tucZ92YR6VmhSva9pJekm2X7nnHWP ?usp=sharing)

```
git clone https://github.com/willy-b/learning-rasp.git
git checkout \
recogs-like-grammar-multiple-templates-supported-but-switch-to-word-level-tokens
python recogs_examples_in_rasp.py
```

The script will show performance on Wu et al 2023 ReCOGS_pos data by default, run with "–use_dev_split", "–use_gen_split" , or "–use_test_split" to see it run on those and give a running score every 10 rows.

For ReCOGS, intending to perform well on Semantic Exact Match, we took a simple, flat, non-tree, non-recursive approach which was able to get approximately 100% on the full test set, and of a sample of 3000 rows of the generalization sets of the real ReCOGS dataset[42]

We took the RASP native sequence tokens, and first did a Transformer learned-embedding compatible operation and created 1 part-of-speech and 4 extra verb-type sequences (because each word in the COGS vocabulary may actually serve multiple POS roles; up to four different verb types as in the case of "liked"

which can serve as v_trans_not_omissible, v_trans_not_omissible_pp_p1, v_trans_not_omissible_pp_p2, and v_cp_taking types).

The five extra sequences serve to associate each word with one or more of the following part-of-speech/verb type roles:

```
det: 1
pp: 2
was: 3
by: 4
to: 5
that: 6
common_noun: 7
proper_noun: 8
v_trans_omissible: 9
v_trans_omissible_pp: 10
v_trans_not_omissible: 11
v_trans_not_omissible_pp: 12
v_cp_taking: 13
v_inf_taking: 14
v_unacc: 15
v_unerg: 16
v_inf: 17
v_dat: 18
v_dat_pp: 19
v_unacc_pp: 20
```

---

[42]word-level token Restricted Access Sequence Processing solution: https://github.com/willy-b/learning-rasp/blob/e97282e18b07004bf714b5c9bb5883090a2ff8e3/word-level-pos-tokens-recogs-style-decoder-loop.rasp

Each of the five sequences comes from a separate map. Since in RASP a map could only have a single value per key, and since individual words had up to four different verb roles (as in "liked" which had 4).

Upon these five parallel, aligned, sequences we used a series of attention head compatible operations to recognize multi-token patterns (see below) corresponding to grammatical forms (listed below).

```
np_det_mask = select(7, pos_tokens, ==)
and select(pos_tokens, 1, ==)
and select(indices+1, indices, ==);
np_prop_mask = select(8, pos_tokens, ==) and
select(indices, indices, ==);
np_det_sequence = aggregate(np_det_mask, 1);
np_prop_sequence = aggregate(np_prop_mask, 1);
np_det_after = select(np_det_sequence, 1, ==) and
select(indices+1, indices, ==);
np_prop_after = select(np_prop_sequence, 1, ==) and
select(indices+1, indices, ==);
np_after_mask = np_det_after or np_prop_after;
np_after_sequence = aggregate(np_after_mask, 1);
np_after_mask = select(np_after_sequence, 1, ==) and
select(indices,indices, ==);
# ...

# np v_unerg
# e.g. [1,7,16]
set example ["the", "guest", "smiled"]
v_unerg_mask = select(16, pos_tokens_vmap1, ==) and
select(indices, indices, ==);
np_v_unerg = aggregate(np_after_mask and v_unerg_mask, 1);
```

These patterns are not causal because their use/input/output is masked to the input section of the sequence, so would take part in the Encoder of the Encoder-Decoder only(all operations outside the input mask in the word-level token RASP solution used in this paper are directly or indirectly causally masked and we built symbol by symbol in a causal autoregressive way). **We could have added an explicit causal mask to each operation but for efficiency and simplicity of the code omitted it when we are doing it implicitly by taking only the last sequence position (we also acausally aggregate so that all sequence positions have the same value as the last sequence position to make it easier to read the output – RASP interpreter will just print it as one position if they are all equal and we only take one position).**

Also, the author thinks many of these RASP steps could be consolidated. The goal here was to first prove by construction that a non-recursive, flat RASP program could get approximately 100% Semantic Exact Match on all the ReCOGS generalization splits (we only missed one split by a little due to two week time constraint, insufficient time to add all rules).

Introduction of variables at the beginning of the ReCOGS logical form (e.g. in the logical form for "a boy painted the girl", we have "boy (1) ; * girl (4) ; painted (2) AND agent (2 , 1) AND theme

(2 , 4)" , the variable introduction is "boy (1) ; * girl (4) ; painted (2)" before the "AND"). A more complete solution that handles not just prepositional phrase recursion (we score approximately 100% with the solution we are describing here) but complement phrases and complement phrase recursion was omitted to fit within the 2 week project, but is believed by the author to be able to be easily done using the same technique (will be done in follow up work). In that approach one would need to loop over the verbs and output the related relationships. Here we simplified and just sorted the input sequence with nouns before verbs and determiners, fillers last (with determiners and fillers not having any corresponding entry in the output sequence). We then count nouns and verbs in the input and count nouns and verbs in the output and determine if we have introduced all the nouns and verbs.

Example counting how many nouns and verbs we have output (introduced as variables) so far (to determine what we need to output for next token):

```
nv_in_output_sequence =
OUTPUT_MASK*(indicator(pos_tokens == 7 or pos_tokens == 8) +
indicator(pos_tokens_vmap1 == 9 or pos_tokens_vmap2 == 10 or
pos_tokens_vmap1 == 11 or pos_tokens_vmap2 == 12 or pos_tokens_vmap3 == 13 or
pos_tokens_vmap4 == 14 or pos_tokens_vmap1 == 15 or pos_tokens_vmap1 == 16 or
pos_tokens_vmap1 == 17 or pos_tokens_vmap1 == 18 or pos_tokens_vmap2 == 19 or
pos_tokens_vmap2 == 20 or pos_tokens_vmap1==21));
nv_in_output_count = selector_width(select(nv_in_output_sequence, 1, ==));
# causal operation as we use only last sequence position
```

How variables are introduced with their index (omitted sorting of input and other operations that can be read in the code and are less important; anything acausal is restricted to input sequence section (Encoder)): (only value at last sequence position is used; causal)

```
# introducing variables
output = "";
# definite article word handling
before_target_word_index = aggregate(select(indices, nv_in_output_count, ==), input_indices_sorted)-1;
has_star = aggregate(select(indices, before_target_word_index, ==), tokens) == "the";
last_output_is_star = aggregate(select(indices, length-1, ==), tokens) == "*";

input_nv_sorted_by_type = input_tokens_sorted_by_type * (input_noun_mask_sorted + input_verb_mask_sorted);
target_word_token = aggregate(select(indices, nv_in_output_count, ==), normalize_nv(input_nv_sorted_by_type))
if (not has_star or last_output_is_star) else "*";
# subtract 1 when matching for producing the index because we just output the additional word by then
target_word_index = aggregate(select(indices, nv_in_output_count-1, ==), input_indices_sorted);

output = target_word_token if ((num_tokens_in_output_excluding_asterisks % 5) == 0) else output;
output = "(" if ((num_tokens_in_output_excluding_asterisks % 5) == 1) else output;
output = target_word_index if ((num_tokens_in_output_excluding_asterisks % 5) == 2) else output;
output = ")" if ((num_tokens_in_output_excluding_asterisks % 5) == 3) else output;
# note that when nv_in_output_count == nv_in_input_count, we will add AND instead of ";"
output = (";" if (5 * nv_in_input_count - 1 > num_tokens_in_output_excluding_asterisks) else "AND")
if (num_tokens_in_output_excluding_asterisks % 5 == 4) else output;

# if we didn't have an input/output separator that needs to be output
output = "|" if num_pipes_in_output == 0 else output;
```

Note that "normalize_nv" is a lookup into a map that has no effect unless the word is a verb in which case it normalizes it to a standard suffix ("ate" to "eat" , "painted" to "paint", etc).

As you can see above, if we have not introduced all the variables, we determine our index into the sorted list of nouns and verbs (nouns before verbs), and using a MLP modeling modulus, compute index mod 5 and condition on that to output that noun/verb or parentheses or index as prediction for next token at last sequence position (all other sequence positions are ignored). Since we do ReCOGS_pos (semantically identical to random indices but avoid requiring random numbers) the index we use is the index of the original noun or verb in the original sequence. If we are still introducing variables, that is the end and we have our prediction for the next token.

If we are done introducing variables at that point in the decoder loop, we move on, and attention head compatible operations recognize templates in the five parallel part-of-speech / verb-type per location sequences for "v_trans_omissible_p1", "v_trans_omissible_p2", "v_trans_omissible_pp_p1", "v_trans_omissible_pp_p2", "v_trans_not_omissible", "v_trans_not_omissible_pp_p1", "v_trans_not_omissible_pp_p2", "v_cp_taking", "v_inf_taking", "v_unacc_p1", "v_unacc_p2", "v_unacc_pp_p1", "v_unacc_pp_p2", "v_unerg", "v_dat_p2", "v_dat_pp_p1", "v_dat_pp_p2", "v_dat_pp_p3", "v_dat_pp_p4".

Here are a couple of examples of patterns, to see how it looks if we support 1 verb pattern per input (no complement phrase recursion; which can be easily handled how we handle other things we loop over, looping over current phrase and masking and processing), which is sufficient to get approximately 100% on all entries that do not use complement phrases (e.g. "so-and-so noticed that (full input here)"):

```
# define the pattern
# ... (just showing one example, np_prop_mask and np_before_mask are attention masks defined earlier)
# np v_dat_p2 np np
# e.g. [8,18,1,7,1,7]
set example ["ella","sold","a","customer","a","car"]
np_np_sequence = aggregate((np_prop_mask and np_before_mask) or (np_det_left_mask and np_two_before_mask), 1);
# would not support prepositional phrase modification on middle NP
#np_np_before_mask = select(np_np_sequence, 1, ==) and select(indices-1, indices, ==);
np_np_any_before_mask = select(np_np_sequence, 1, ==) and select(indices, indices, >); # acausal is ok in INPUT sequence (encoder part, not decode
np_np_any_before_sequence = aggregate(np_np_any_before_mask, 1);
np_np_any_before_mask = select(np_np_any_before_sequence, 1, ==) and select(indices, indices, ==);
np_v_dat_p_np_np = aggregate(np_after_mask and v_dat_mask and np_np_any_before_mask, 1);
# Example: np_v_dat_p_np_np(['ella', 'sold', 'a', 'customer', 'a', 'car']) = [0, 1, 0, 0, 0, 0] (ints)
# Example: np_v_dat_p_np_np([8, 18, 1, 7, 1, 7]) = [0, 1, 0, 0, 0, 0] (ints)

# ...

# check the pattern and set the template name
any_np_v_trans_omissible = aggregate(select(np_v_trans_omissible, 1, ==), 1);
template_name = "v_trans_omissible_p1"
if (any_np_v_trans_omissible == 1) else template_name;

# ...

any_v_dat_p2 = aggregate(select(np_v_dat_p_np_np, 1, ==), 1);
template_name = "v_dat_p2" if (any_v_dat_p2 == 1) else template_name;

# ...

any_v_dat_pp_p4 = aggregate(select(np_was_v_dat_pp_np_by_np, 1, ==), 1);
```

```
template_name = "v_dat_pp_p4" if (any_v_dat_pp_p4 == 1) else template_name;

# must be checked after P4
any_v_dat_pp_p2 = aggregate(select(np_was_v_dat_pp_to_np_by_np, 1, ==), 1);
template_name = "v_dat_pp_p2" if (any_v_dat_pp_p2 == 1) else template_name;

# template name is used to lookup the number of verb relationships to output and
```

The rest of this applies to just values used from the last sequence location (output is prediction for next symbol).

Based on the template recognized, we lookup the template size for number of relationships (theme, recipient, agent) for that verb type:

```
def template_size(template_name) {
  template_sizes = {
 "": 0,
 "v_trans_omissible_p1": 1,
 "v_trans_omissible_p2": 2,
 "v_trans_omissible_pp_p1": 1,
 "v_trans_omissible_pp_p2": 2,
 "v_trans_not_omissible": 2,
 "v_trans_not_omissible_pp_p1": 1,
 "v_trans_not_omissible_pp_p2": 2,
 "v_cp_taking": 2,
 "v_inf_taking": 4,
 "v_unacc_p1": 2,
 "v_unacc_p2": 1,
 "v_unacc_pp_p1": 1,
 "v_unacc_pp_p2": 2,
 "v_unerg": 1,
# "v_inf": 1,
 "v_dat_p1": 3,
 "v_dat_p2": 3,
 "v_dat_pp_p1": 2,
 "v_dat_pp_p2": 3,
 "v_dat_pp_p3": 2,
 "v_dat_pp_p4": 3
  };
  # v_inf_taking includes v_inf and an extra verb is why it is 4 instead of 2
  return template_sizes[template_name];
}
```

Details are in the code, but we compute at the last sequence position (in parallel) the number of relationships output for the verb so far, and for the current relationship which token within that multi-token process (e.g. the word "agent" or the open parenthesis "(" or the left index, or the comma, or right index, close parenthesis ")", "AND", etc) we are on.

Like we computed at the last sequence position the number of nouns and verbs in the output once we are finished introducing nouns and verbs (this would be a little different with complement phrases[43], we compute the number of agent,theme,recipient,xcomp,ccomp entries in the output:

```
atrxc_in_output_sequence = OUTPUT_MASK*(indicator(tokens == "agent"
or tokens == "theme"
or tokens=="recipient"
or tokens=="xcomp" or tokens =="ccomp"));
# agent_theme_recipient_xcomp_ccomp_output_count is the number of relationships we have output
agent_theme_recipient_xcomp_ccomp_output_count =
selector_width(select(atrxc_in_output_sequence, 1, ==));
after_intro_idx =
(nv_in_output_count - nv_in_input_count + agent_theme_recipient_xcomp_ccomp_output_count) if nv_in_output_count >= nv_in_input_count else 0;
after_intro_num_tokens_in_output_excluding_asterisks =
num_tokens_in_output_excluding_asterisks - ((5 * nv_in_input_count));
```

---

[43]see Complement phrase support work completed in https://github.com/willy-b/learning-rasp/pull/7

We use all those different values which are computed independently from the input sequence and so do not add depth/layer requirements as many of the ones which involve accessing the sequence can be done in parallel. We then use the verb-type and relationship index to the relationship into a map to get the current relationship to output (as some verb types output agent first, some output theme, etc):

```
template_mapping1 = {
 "": "",
 "v_trans_omissible_p1": "agent",
 "v_trans_omissible_p2": "agent",
 "v_trans_omissible_pp_p1": "theme",
 "v_trans_omissible_pp_p2": "theme",
 "v_trans_not_omissible": "agent",
 "v_trans_not_omissible_pp_p1": "theme",
 "v_trans_not_omissible_pp_p2": "theme",
 "v_cp_taking": "agent",
 "v_inf_taking": "agent",
 "v_unacc_p1": "agent",
 "v_unacc_p2": "theme",
 "v_unacc_pp_p1": "theme",
 "v_unacc_pp_p2": "theme",
 "v_unerg": "agent",
 "v_inf": "agent",
 "v_dat_p1": "agent",
 "v_dat_p2": "agent",
 "v_dat_pp_p1": "theme",
 "v_dat_pp_p2": "theme",
 "v_dat_pp_p3": "recipient",
 "v_dat_pp_p4": "recipient"
};
```

Outputting the verb relationships we must skip over any "pp np" as possible agents, themes, or recipients to avoid getting confused by noun phrases added by prepositional modification (believed by the author to be the cause of the issue with obj pp to subj pp generalization by (Wu et al., 2024)'s Transformer).

```
pp_sequence = indicator(pos_tokens == 2);
pp_one_after_mask = select(pp_sequence, 1, ==) and select(indices+1, indices, ==);
pp_one_after_sequence = aggregate(pp_one_after_mask, 1);
pp_one_after_mask = select(pp_one_after_sequence, 1, ==) and select(indices, indices, ==);

pp_two_after_mask = select(pp_sequence, 1, ==) and select(indices+2, indices, ==);
pp_two_after_sequence = aggregate(pp_two_after_mask, 1);
pp_two_after_mask = select(pp_two_after_sequence, 1, ==) and select(indices, indices, ==);

np_det_diag_mask = select(aggregate(np_det_mask, 1), 1, ==) and select(indices, indices, ==);
np_prop_diag_mask = select(aggregate(np_prop_mask, 1), 1, ==) and select(indices, indices, ==);

no_pp_np_mask =
1 - aggregate((pp_one_after_mask and np_prop_diag_mask) or
(pp_two_after_mask and np_det_diag_mask), 1);
nps_without_pp_prefix_indices = selector_width(select(NOUN_MASK*no_pp_np_mask, 1, ==) and select(indices, indices, <=))*NOUN_MASK*no_pp_np_mask;

left_idx = aggregate(select(indices, left_idx_in_nps_zero_based, ==), input_indices_sorted);
right_idx = aggregate(select(nps_without_pp_prefix_indices, after_intro_idx, ==), indices);

# points to 2nd verb for xcomp for v_inf_taking_v_inf
right_idx = aggregate(select(indices, (nv_in_output_count-1), ==), input_indices_sorted)
if (template_name == "v_inf_taking" and after_intro_idx == 2) else right_idx;

# points to 1st noun for 2nd v_inf agent in v_inf_taking_v_inf
right_idx = aggregate(select(indices, 0, ==), input_indices_sorted)
if (template_name == "v_inf_taking" and after_intro_idx == 3) else right_idx;

# ...

after_intro_target_token = left_idx
if (after_intro_num_tokens_in_output_excluding_asterisks % 7 == 2) else after_intro_target_token;

after_intro_target_token = ","
if (after_intro_num_tokens_in_output_excluding_asterisks % 7 == 3) else after_intro_target_token;

after_intro_target_token = right_idx
if (after_intro_num_tokens_in_output_excluding_asterisks % 7 == 4)
else after_intro_target_token;

after_intro_target_token = ")"
if (after_intro_num_tokens_in_output_excluding_asterisks % 7 == 5)
else after_intro_target_token;

after_intro_target_token = "AND"
if (after_intro_num_tokens_in_output_excluding_asterisks % 7 == 6
and not (template_mapping_output == "")) else after_intro_target_token;

# ...
```

After outputting all verb relationships, we consider whether we have prepositional phrase noun modifiers to record in the logical form.

For outputting prepositional relationships ("nmod . in", "nmod . on", "nmod . beside"), we do a similar approach, counting prepositional phrases in the input, counting how many nmods we have output, figuring out which one is currently being output:

```
pps_in_input_sequence = INPUT_MASK*(indicator(pos_tokens == 2));
pps_in_input_count = selector_width(select(pps_in_input_sequence, 1, ==));
pps_index = pps_in_input_sequence*selector_width(select(pps_in_input_sequence, 1, ==)
and select(indices,indices, <=));
nmods_and_pps_in_output_sequence = OUTPUT_MASK*(indicator(tokens == "nmod . in" or tokens == "nmod . beside" or tokens=="nmod . on"));
nmods_and_pps_in_output_count = selector_width(select(nmods_and_pps_in_output_sequence, 1, ==));

current_pp = aggregate(select(pps_index, nmods_and_pps_in_output_count+1, ==), tokens) if pps_in_input_count > 0 else "";
current_pp = "" if current_pp == 0 else current_pp;
current_nmod_token =
("nmod . " + current_pp) if (pps_in_input_count > 0 and not (current_pp == 0)
and after_intro_num_tokens_in_output_excluding_asterisks % 7 == 0) else "";
current_nmod_token = "(" if after_intro_num_tokens_in_output_excluding_asterisks % 7 == 1 else current_nmod_token;
current_nmod_token =
(aggregate(select(pps_index, nmods_and_pps_in_output_count, ==), indices)-1) if pps_in_input_count > 0
and after_intro_num_tokens_in_output_excluding_asterisks % 7 == 2 else current_nmod_token;
current_nmod_token = ","
if after_intro_num_tokens_in_output_excluding_asterisks % 7 == 3 else current_nmod_token;
after_nmod_idx =
aggregate(select(pps_index, nmods_and_pps_in_output_count, ==), indices)+1;
token_at_after_nmod_idx =
aggregate(select(indices, after_nmod_idx, ==), tokens);
after_nmod_idx = (after_nmod_idx + 1) if (token_at_after_nmod_idx == "the" or token_at_after_nmod_idx == "a") else after_nmod_idx;
current_nmod_token = (after_nmod_idx)
if pps_in_input_count > 0
and after_intro_num_tokens_in_output_excluding_asterisks % 7 == 4 else current_nmod_token;
current_nmod_token = ")"
if after_intro_num_tokens_in_output_excluding_asterisks % 7 == 5
else current_nmod_token;
current_nmod_token =
("AND" if nmods_and_pps_in_output_count < pps_in_input_count else "")
if after_intro_num_tokens_in_output_excluding_asterisks % 7 == 6
else current_nmod_token;
after_intro_and_relationships_nmod_token =
current_nmod_token if nmods_and_pps_in_output_count <= pps_in_input_count else "";
num_tokens_in_nmod_section =
after_intro_num_tokens_in_output_excluding_asterisks - template_size(template_name)*7 + 1;
```

See code for full details. For all steps only the RASP outputs aligned with the input sequence (Encoder part of derived Transformer) or the very last sequence output (for next token in autoregressive generation) were used. For convenience of reading the aggregate operator was usually used acausally to assign all sequence outputs before the last one to the same value as the last (so only one value would be displayed).

You can run a demo and see the autoregressive output

(or just visit https://colab.research.google.com/drive/ 1FS4tucZ92YR6VmhSva9pJekm2X7nnHWP ?usp=sharing)

```
git clone https://github.com/willy-b/learning-rasp.git
git checkout \
recogs-like-grammar-multiple-templates-supported-but-switch-to-word-level-tokens
python recogs_examples_in_rasp.py
```

18 Appendix 6 - Note on a Restricted Access Sequence Processing character-level token program / model design (NOT what is used in this paper but feasible)

Note, a proof of concept character level Restricted Access Sequence Processing model was done with a decoder loop (unlike word-level solution above, it was a sketch so did not limit to strictly causal operations which just require more careful indexing – using the value at the separator or the end of a word instead of pooling the same value to all letters in a word for example). Note that this one did not cover translating sentences in general into ReCOGS unlike the word-level solution as it is tedious and redundant but the core operations are possible and the author believes any solution at the word level can be mapped to a solution in character level tokens (out of scope for this paper to prove it).

Since it is a separate problem and adds a lot of complexity without bringing anything to bear on the main questions of the paper, I left a full implementation to the word-level tokens which were simpler and ran faster. The difference is one uses a similar approach started at [44] to assign all the letters in each word an index.

Word indices can be assigned using RASP to count separators occuring prior to each sequence location like:

(we also zero out the word index for the separators themselves)

```
word_indices = \
(1+selector_width(select(tokens, " ", ==)
and select(indices, indices, <=)))
*(0 if indicator(tokens == " ") else 1);
```

Then one can do an aggregation of the letters grouping by word index (this, which is NOT part of the techniques used in this paper for the word-level tokens solution, requires additional work (tedious not challenging) to do causally outside the input (in the decoder), one must sum forward so the word representation is always at the last letter of the word or separator instead of at all letters of the word, and that step is left out of the character-level demo and this discussion; whereas the word-level solution described above has a clear Encoder Decoder separation. This can be done so that the value which is then the same for all letters in each word,

is unique to each word in the dictionary and can be looked up in a map to get word level attributes like part-of-speech and get back to the solution in the word-level tokens in Appendix 5 which was fully implemented. A simple approach (not necessarily recommended but works for proof of concept) that would work for small vocabularies (easily extended) is to use a map to lookup each letter of the alphabet to a log prime. Then the sum of the letters in a word (grouped by the word index which is the count of spaces/separators prior) is the sum of the log primes indexed by the alphabet index. Since the sum of logarithms of numbers is the same as the logarithm of the product of those numbers, this is equivalent to the logarithm of the product of a series of primes. Each prime in the product corresponds 1-to-1 to a letter in the alphabet, with the number of occurrences in the product corresponding to the number of times that letter occurs in the word. By uniqueness of prime number factorization this would map each multiset of letters to a single unique sum of log primes. Thus if you do not have words which are anagrams, all the letters in each word would be assigned a number that uniquely represented that word in the vocabulary. If you have anagrams you can do this step and then take the first and last letter and compute a separate number from that and add it to all the letters in the word.

Example lookup table for letters before aggregating by word index (not recommended but for proof of concept that one can go from character level tokens to word-specific numbers which can then be looked up as in the word-level token solution in Appendix 5 used throughout the paper):

---

[44]https://github.com/willy-b/learning-rasp/blob/main/decoder-loop-example-parse-into-recogs-style-variables.rasp

```
def as_num_for_letter_multiset_word_pooling(t) {
    # To be multiset unique, need logarithm of prime so that the sum aggregation
    # used in RASP corresponds to prime number factorization (sum of logs of primes is same as log of product of primes)
    # (we can do sum aggregation instead of mean by multiplying by length)
    # However RASP does not appear to support logarithms (underlying multilayer
    # perceptron can learn to approximate logarithms)
    #letter_to_prime_for_multiset_word_pooling = {"a": 2, "b": 3, "c": 5, "d": 7,
    #"e": 11, "f": 13, "g": 17, "h": 19, "i": 23, "j": 29, "k": 31, "l": 37,
    #"m": 41, "n": 43, "o": 47, "p": 53, "q": 59, "r": 61, "s": 67, "t": 71,
    #"u": 73, "v": 79, "w": 83, "x": 89, "y": 97, "z": 101, ".": 0,
    #" ": 0, ":": 0};
    map_letter_to_log_prime_for_pooling = {"a": 0.6931471805599453, "b": 1.0986122886681098,
    "c": 1.6094379124341003, "d": 1.9459101490553132, "e": 2.3978952727983707,
    "f": 2.5649493574615367, "g": 2.833213344056216, "h": 2.9444389791664403,
    "i": 3.1354942159291497, "j": 3.367295829986474, "k": 3.4339872044851463,
    "l": 3.6109179126442243, "m": 3.713572066704308, "n": 3.7612001156935624,
    "o": 3.8501476017100584, "p": 3.970291913552122, "q": 4.07753744390572,
    "r": 4.110873864173311, "s": 4.204692619390966, "t": 4.2626798770413155,
    "u": 4.290459441148391, "v": 4.3694478524670215, "w": 4.418840607796598,
    "x": 4.48863636973214, "y": 4.574710978503383, "z": 4.61512051684126,
    # we zero out tokens we want not to affect the identity of the word
    ".": 0, " ": 0, ":": -1, "(": -1, ")": -1, "0": -1, "1": -1, "2": -1,
    "3": -1, "4": -1, "5": -1, "6": -1, "7": -1, "8": -1, "9": -1, ";": -1,
    ",": -1};
    return map_letter_to_log_prime_for_pooling[t];
}
```

Pooling by word can then be done with:

```
pseudoembeddedwords = aggregate(select(word_indices, word_indices, ==), as_num_for_letter_multiset_word_pooling(tokens))*word_lengths;
```

(Per-character token example is not causally masked, we do causal strict-decoder-compatible solution for anything outside input sequence in the full word-level solution above just leaving out of this character-level sketch, which is NOT used in this paper. For the causal character level solution one would use the summed value at the end of the word or the separator instead, indexing relative to separators.)

Those values could then be looked up in a dictionary like in the completed word-level token solution to get part-of-speech, verb-type, etc, to derive a separate sequence which can be used for template matching as we successfully did with word-level tokens (see Appendix 5).

19   Appendix 7 - Computing Grammar Coverage

First we use the grammar as it was generated as a probablistic context free grammar per (Kim and Linzen, 2020) using the full details put in Lark format by (Klinger et al., 2024) and converting it ourselves to a format compatible with (Zeller et al., 2023).

Note this starting point is not the grammar we claim the our Restricted Access Sequence Processing model implements or the Transformer actually learns as we argue the Transformer is learning a flat, non-tree solution to this simple grammar (not actually learning to collapse "np_det pp np" into "np" for example). First we compute grammar coverage relative to the PCFG approach that generated it, which mostly aligns with our RASP model. We also ignore terminals in this assessment of coverage, as stated earlier, when computing grammar coverage, we will report the grammar coverage over expansions that collapse all vocabulary leaves to a single leaf (for example not requiring that every particular proper noun or common noun be observed in a particular pattern, so long as one has and we can confirm the code treats them as equivalent; e.g. having tested "Liam drew the cat" and proven that "Liam" and "Noah" are treated as interchangeable proper nouns, and that "cat" and "dog" are treated as interchangeable common nouns by the RASP solution – not something one can assume for neural network solutions in general – means that confirming our solution produces the correct logical form for "Liam drew the cat" suffices to prove the RASP solution can handle "Noah drew the dog", which saves us a lot of work so long as we make sure to write our RASP solution such that noah/liam and cat/dog are indeed treated identically).

```
COGS_INPUT_GRAMMAR_NO_TERMINALS = {
"<start>": ["<s1>", "<s2>", "<s3>", "<s4>", "<vp_internal>"],
"<s1>": ["<np> <vp_external>"],
"<s2>": ["<np> <vp_passive>"],
"<s3>": ["<np> <vp_passive_dat>"],
"<s4>": ["<np> <vp_external4>"],
"<vp_external>": ["<v_unerg>", "<v_trans_omissible_p1>", "<vp_external1>", "<vp_external2>", "<vp_external3>", "<vp_external4>","<vp_external5>",
"<vp_external1>": ["<v_unacc_p1> <np>"],
"<vp_external2>": ["<v_trans_omissible_p2> <np>"],
"<vp_external3>": ["<v_trans_not_omissible> <np>"],
"<vp_external4>": ["<v_inf_taking> <to> <v_inf>"],
"<vp_external5>": ["<v_cp_taking> <that> <start>"],
```

```
"<vp_external6>": ["<v_dat_p1> <np> <pp_iobj>"],
"<vp_external7>": ["<v_dat_p2> <np> <np>"],
"<vp_internal>": ["<np> <v_unacc_p2>"],
"<vp_passive>": ["<vp_passive1>", "<vp_passive2>", "<vp_passive3>", "<vp_passive4
"<vp_passive1>": ["<was> <v_trans_not_omissible_pp_p1>"],
 "<vp_passive2>": ["<was> <v_trans_not_omissible_pp_p2> <by> <np>"] ,
"<vp_passive3>": ["<was> <v_trans_omissible_pp_p1>"],
"<vp_passive4>": ["<was> <v_trans_omissible_pp_p2> <by> <np>"],
"<vp_passive5>": ["<was> <v_unacc_pp_p1>"],
"<vp_passive6>": ["<was> <v_unacc_pp_p2> <by> <np>"],
"<vp_passive7>": ["<was> <v_dat_pp_p1> <pp_iobj>"],
"<vp_passive8>": ["<was> <v_dat_pp_p2> <pp_iobj> <by> <np>"],
"<vp_passive_dat>": ["<vp_passive_dat1>", "<vp_passive_dat2>"],
"<vp_passive_dat1>": ["<was> <v_dat_pp_p3> <np>"],
"<vp_passive_dat2>": ["<was> <v_dat_pp_p4> <np> <by> <np>"],
"<np>": ["<np_prop>","<np_det>", "<np_pp>"],
"<np_prop>": ["<proper_noun>"],
"<np_det>": ["<det> <common_noun>"],
"<np_pp>": ["<np_det> <pp> <np>"],
"<pp_iobj>": ["<to> <np>"],
"<det>": [],
"<pp>": [],
"<was>": [],
"<by>": [],
"<to>": [],
"<that>": [],
"<common_noun>": [],
"<proper_noun>": [],
"<v_trans_omissible_p1>": [],
"<v_trans_omissible_p2>": [],
"<v_trans_omissible_pp_p1>": [],
"<v_trans_omissible_pp_p2>": [],
"<v_trans_not_omissible>": [],
"<v_trans_not_omissible_pp_p1>": [],
"<v_trans_not_omissible_pp_p2>": [],
"<v_cp_taking>": [],
"<v_inf_taking>": [],
"<v_unacc_p1>": [],
"<v_unacc_p2>": [],
"<v_unacc_pp_p1>": [],
"<v_unacc_pp_p2>": [],
"<v_unerg>": [],
"<v_inf>": [],
"<v_dat_p1>": [],
"<v_dat_p2>": [],
"<v_dat_pp_p1>": [],
"<v_dat_pp_p2>": [],
"<v_dat_pp_p3>": [],
"<v_dat_pp_p4>": [],
}
```

After parsing a sentence with the Lark parser, we can compute the expansions it covers with the following Python:

```python
def generate_set_of_expansion_keys_for_lark_parse_tree(tree):
  nodes = [tree]
  expansions_observed = set()
  for node in nodes:
    current_node_label = node.data[:]
    children = node.children
    expansion = f"<{current_node_label} ->"
    for child in children:
      # add expansion for current -> child
      child_node_label = child.data[:]
      expansion += f" <{child_node_label}>"
      # also process expansions from child
      nodes.append(child)
    if len(children) > 0:
      #print(f"{expansion}")
      expansions_observed.add(expansion)
  return expansions_observed
```

For example, for the sentence "the girl noticed that a boy painted the girl",

(note we otherwise exclude such complement phrases in this paper[45], which you can see as " <v_cp_taking> <that>" below)

we get

```
sentence = "the girl noticed that a boy painted the girl"
tree = parser.parse(sentence)
expansions_observed = generate_set_of_expansion_keys_for_lark_parse_tree(tree)
# <start> -> <s1>
# <s1> -> <np> <vp_external>
```

---

[45]Complement phrase support has been added in https://github.com/willy-b/learning-rasp/pull/7 , and the model scores 100% semantic exact match (95% confidence interval: 99.88% to 100.00%) on the full ReCOGS_pos Test set, and 100% exact string match (same confidence interval) on the dev set (test set results pending).

```
# <np> -> <np_det>
# <vp_external> -> <vp_external5>
# <np_det> -> <det> <common_noun>
# <vp_external5> -> <v_cp_taking> <that> <start>
# <start> -> <s1>
# <s1> -> <np> <vp_external>
# <np> -> <np_det>
# <vp_external> -> <vp_external2>
# <np_det> -> <det> <common_noun>
# <vp_external2> -> <v_trans_omissible_p2> <np>
# <np> -> <np_det>
# <np_det> -> <det> <common_noun>
```

At first we use TrackingGrammarCoverage-Fuzzer (from (Zeller et al., 2023)) to compute the set of all possible grammar expansions:

```
cogs_simplified_input_grammar_fuzzer = TrackingGrammarCoverageFuzzer(COGS_INPUT_GRAMMAR_SIMPLIFIED)
expected_expansions = cogs_simplified_input_grammar_fuzzer.max_expansion_coverage()
```

One can use this to get a sense of what it is possible to learn about the grammar from a particular set of examples

and what examples need to be seen at a minimum for any model to learn the task from scratch and could possibly help one design a minimum length dataset with low redundancy. Note for a Transformer model learning word embeddings / mapping to part-of-speech for each word, one would need to use the grammar with terminals to compute coverage. Here we want to argue something about our RASP model where we can ensure via implementation that all terminals in a category are treated identically (and we observe 100% semantic exact match for the related generalization splits for swapping words within a part of speech).

We can ask what % of the grammar without terminals is covered by the first 20 sentences from COGS?

```
nonsense_example_sentences = [
"A rose was helped by a dog",
"The sailor dusted a boy",
"Emma rolled a teacher",
"Evelyn rolled the girl",
"A cake was forwarded to Levi by Charlotte",
"The captain ate",
"The girl needed to cook",
"A cake rolled",
"The cookie was passed to Emma",
"Emma ate the ring beside a bed",
"A horse gave the cake beside a table to the mouse",
"Amelia gave Emma a strawberry",
"A cat disintegrated a girl",
"Eleanor sold Evelyn the cake",
"The book was lended to Benjamin by a cat",
"The cake was frozen by the giraffe",
"The donut was studied",
"Isabella forwarded a box on a tree to Emma",
"A cake was stabbed by Scarlett",
"A pencil was fed to Liam by the deer",
]

all_expansions_observed_across_examples = set()

for sentence in nonsense_example_sentences:
  single_example_expansions = generate_set_of_expansion_keys_for_lark_parse_tree(parser.parse(sentence.lower()))
  all_expansions_observed_across_examples = all_expansions_observed_across_examples.union(single_example_expansions)

1 - len(set(expansions_expected) - all_expansions_observed_across_examples) / len(expansions_expected)
# 0.6603773584905661
```

Those 20 sentences cover 66% of the grammar. (Continued on next page.)

We can compare the first 20 sentences of COGS that to the 19 sentences used in developing the RASP program plus one to cover basic preposi-

tional phrases:[46]

```
handpicked_example_sentences = [ # missing complement phrase (cp_taking) examples
"the girl was painted", # see link above all these examples each correspond to di
"a boy painted",
"a boy painted the girl",
"the girl was painted by a boy",
"a boy respected the girl",
"the girl was respected",
"the girl was respected by a boy",
"the boy grew the flower",
"the flower was grown",
"the flower was grown by a boy",
"the scientist wanted to read",
"the guest smiled",
"the flower grew",
"ella sold a car to the customer",
"ella sold a customer a car",
"the customer was sold a car",
"the customer was sold a car by ella",
"the car was sold to the customer by ella",
"the car was sold to the customer",
]

all_expansions_observed_across_examples = set()

for sentence in handpicked_example_sentences:
  single_example_expansions = generate_set_of_expansion_keys_for_lark_parse_tree(
  all_expansions_observed_across_examples = all_expansions_observed_across_exampl

1 - len(set(expansions_expected) - all_expansions_observed_across_examples) / len
# 0.9056603773584906

set(expansions_expected) - all_expansions_observed_across_examples
#{'<np> -> <np_pp>', # tells us we need a prepositional phrase example!
# '<np_pp> -> <np_det> <pp> <np>', # tell us we need prepositional phrase example
# '<vp_external5> -> <v_cp_taking> <that> <start>',# we know we skipped complemen
# '<vp_external> -> <vp_external4>', # appears redundant path to get here in gram
# '<vp_external> -> <vp_external5>'} # complement phrase related, skipped in this
```

We got to 90.5% grammar coverage in our 19 examples (higher if we ignore the 2 CP related entries and redundant vp_external4) instead of COGS 66% in 20 examples.

And, it is telling us we are missing an example with prepositional phrases (see next example).

Let us add a simple prepositional phrase example:

```
handpicked_example_sentences = \
handpicked_example_sentences + \
["a boy painted the girl in a house"]

handpicked_example_sentences
# ['the girl was painted',
#  'a boy painted',
#  'a boy painted the girl',
#  'the girl was painted by a boy',
#  'a boy respected the girl',
#  'the girl was respected',
#  'the girl was respected by a boy',
#  'the boy grew the flower',
#  'the flower was grown',
#  'the flower was grown by a boy',
#  'the scientist wanted to read',
#  'the guest smiled',
#  'the flower grew',
#  'ella sold a car to the customer',
#  'ella sold a customer a car',
#  'the customer was sold a car',
#  'the customer was sold a car by ella',
#  'the car was sold to the customer by ella',
#  'the car was sold to the customer',
#  'a boy painted the girl in a house']
all_expansions_observed_across_examples = set()

for sentence in handpicked_example_sentences:
  single_example_expansions = generate_set_of_expansion_keys_for_lark_parse_tree(parser.parse(sentence.lower()))
  all_expansions_observed_across_examples = all_expansions_observed_across_examples.union(single_example_expansions)

1 - len(set(expansions_expected) - all_expansions_observed_across_examples) / len(expansions_expected)
# 0.9433962264150944

set(expansions_expected) - all_expansions_observed_across_examples
#{'<vp_external5> -> <v_cp_taking> <that> <start>', # complement phrase related
# '<vp_external> -> <vp_external4>', # redundant
# '<vp_external> -> <vp_external5>'} # complement phrase related
```

(continued below)

Thus we saw what we did not cover yet and in 19 intentionally crafted sentences (each is in the RASP code with a corresponding rule) cover 90% of the grammar, we add one sentence to fill the reported gap and get to 94% and possibly 100% of what we need (I wrote the RASP program primarily on these examples derived from the grammar with complement phrases excluded and after including minimal prepositional phrase handling (by skipping "pp np" when choosing nouns for relationships and setting up an nmod output loop) I indeed scored 95% on the first hundred examples from the training set, after which I made small modifications and got >99% on the test and gen sets – I knew there were prepositional patterns not yet covered but it is a two week project and time to measure).

However these coverage metrics are misleading when it comes to prepositional phrases as it would not suggest to include prepositional phrases in all positions, assuming they could be collapsed by the model back to 'np' using 'np -> np_pp -> np_det pp np' while experiments suggest it is necessary to train with prepositional phrases explicitly in the different positions in the different verb patterns (see below).

Based on the finding earlier we believe that the only recursion learned is tail recursion in the decoder loop and that 'np -> np_det | np_prop | np_pp' and 'np_pp -> np_det pp np' is not actually performed as if the Encoder-Decoder Transformer were to learn a tree-based or recursive representation. If the Transformer had a tree based representation, it is predicted that the "v_dat_p2_pp_moved_to_recipient" would not be any harder than when the pp modification is on the theme, as 'np v_dat_p2 np_det pp np np' can be transformed by the recursive grammar rule 'np_det pp np -> np_pp -> np' to 'np v_dat_p2 np np' on which it is already trained and has good performance, but instead it fails completely, and see also "Error Analysis for Wu et al 2023 baseline Encoder-Decoder Transformer on obj_pp_to_subj_pp split" where we observe that prepositional modification of a noun to the left of a verb it is the agent of causes the new prepositional phrase noun that becomes the closest noun to be mistaken for the agent, which is in contradiction to the model collapsing 'np_det pp np' to 'np' before matching the overall grammar pattern.

That said with a couple of simple rules that were not tree we were able to get 100% on the

pp_recursion split (up to depth 12) and approximately 90% of the obj_pp_to_subj_pp split.

Modifying the grammar coverage to model this non-tree representation would be exciting to address in future work.

20   Appendix 8 - Model Detail

For our Restricted Access Sequence Processing ReCOGS program, we used the RASP interpreter of (Weiss et al., 2021) to run our program. For RASP model design and details see Appendix 5. We use word-level tokens for all results in this paper.[47] Consistent with (Zhou et al., 2023) we use (Weiss et al., 2021)'s RASP originally used for modeling Transformer encoders to model an encoder-decoder in a causal way by feeding the autoregressive output back into the program. We only have aggregations with non-causal masks when that aggregation (or without loss of generality just before the aggregation product is used to avoid multiplying everywhere) is masked by an input mask restricting it to the sequence corresponding to the input.[48]

We used RASP maps to map word level tokens to part-of-speech and verb-type which is consistent with what can be learned in embeddings or the earliest layer of a Transformer (Tenney et al., 2019) and then did 19 different attention-head based template matches on a flat sequence[49] (no tree-based parsing, no recursive combination of terminals/nonterminals) to match the sentence to a template in the grammar (see "Appendix 5 - Restricted Access Sequence Processing word-level (post-embedding) token program/model design" and "Appendix 1: Vocabulary and Grammar"). The 19 types were based on (Zeller et al., 2023) grammar coverage of the COGS grammar (see Methods and "Appendix 7 - Computing Grammar Coverage").[50]

---

[47]We believe any solution at the word-level can be converted to a character-level token solution and that is not the focus of our investigation here (see Appendix 6 for proof of concept details on a character level solution not used here).

[48]An example the author has prepared of this is available at https://github.com/willy-b/learning-rasp/blob/16a8e154b025e91c8e56965a1d475e49f69ebdbd/recogs_examples_in_rasp.py .

[49]A flat/non-tree solution was pursued because it was simple and given the failure documented in (Wu et al., 2024) of the baseline Encoder-Decoder to generalize from obj_pp_to_subj_pp and other evidence we give below we shall see it is hard to argue a tree-based solution which includes the rule 'np_det pp np -> np_pp -> np' is learned by (Wu et al., 2024)'s baseline Encoder-Decoder Transformer.

[50]To handle prepositional phrases in a flat solution, we find it necessary to add a rule that ignores noun phrases preceded by a prepositional phrase (ignore "pp np") when searching for

We also ensure that the mapping from words to part-of-speech and verb type is complete based on a published list of such mappings and put that into our hardcoded word embedding[51]

For training Transformers from scratch with randomly initialized weights using gradient descent for comparison with RASP predictions, we use scripts derived from those provided by (Wu et al., 2024)[52].

21 Appendix 9 - Methods Detail

We use the RASP (Weiss et al., 2021) intepreter[53] to evaluate our RASP programs[54].

We implement in RASP the transformation of COGS input sentences into ReCOGS_pos[55]. logical forms (LFs) which are scored by Semantic Exact Match[56] against ground truth.

In the training data only, any ReCOGS training augmentations like preposing or "um" sprinkles are excluded when evaluating the RASP model on the train data (it does not learn directly from the examples and these augmentations are outside of the grammar).

For the RASP solution, due to time constraints, complement phrase examples e.g. "Emma noticed that ..." which it appears can be implemented in RASP, were excluded so the author could focus on

---

noun indexes to report in relationships (agent, theme, recipient, etc), and we loosen verb type templates to allow a gap for any prepositional phrase to be inserted (this was not done in every possible case manually within 2 weeks which is why we do not get 100% in that one split, see next results section).

[51] It is reported that pretrained transformers seem to have learned POS information at their earliest layers, (Tenney et al., 2019)

[52] https://github.com/frankaging/ReCOGS/blob/1b6eca8ff4dca5fd2fb284a7d470998af5083beb/run_cogs.py
and
https://github.com/frankaging/ReCOGS/blob/1b6eca8ff4dca5fd2fb284a7d470998af5083beb/model/encoder_decoder_hf.py

[53] provided at https://github.com/tech-srl/RASP/

[54] https://github.com/willy-b/learning-rasp/blob/16a8e154b025e91c8e56965a1d475e49f69ebdbd/word-level-pos-tokens-recogs-style-decoder-loop.rasp
with a demo at
https://github.com/willy-b/learning-rasp/blob/16a8e154b025e91c8e56965a1d475e49f69ebdbd/recogs_examples_in_rasp.py

[55] We use the ReCOGS positional index data (rather than default ReCOGS with randomized indices) as it has consistent position based indices that allow us to perform well on Exact Match (like the original COGS task) as well as Semantic Exact Match (which ignores absolute values of indices).
See ReCOGS_pos dataset at
https://github.com/frankaging/ReCOGS/tree/1b6eca8ff4dca5fd2fb284a7d470998af5083beb
/recogs_positional_index

[56] https://github.com/frankaging/ReCOGS/blob/1b6eca8ff4dca5fd2fb284a7d470998af5083beb/utils/train_utils.py
and
https://github.com/frankaging/ReCOGS/blob/1b6eca8ff4dca5fd2fb284a7d470998af5083beb/utils/compgen.py

---

full coverage of the remaining grammar within the two weeks.[57].

We also measure grammar coverage of input examples supported by our RASP model against the full grammar of COGS/ReCOGS input sentences provided in the utilities of the IBM CPG project (Klinger et al., 2024)[58]

When computing grammar coverage (Zeller et al., 2023), we collapse all vocabulary terminals (leaves) to a single terminal (leaf), ignoring purely lexical differences (see Appendix 7 for details and motivation).

The overall Semantic Exact Match performance is reported as well as the performance on the specific structural generalization splits where Transformers are reported to struggle, even in ReCOGS, specifically Object Prepositional Phrase to Subject Prepositional Phrase (obj_pp_to_subj_pp), Prepositional Phrase (pp_recursion) will be highlighted and discussed in depth for all models.

For the RASP program's Semantic Exact Match results which is based on the outcome of a deterministic program (so cannot randomly reinitialize weights and retrain, rerun), we can use the Beta distribution to model the uncertainty and generate confidence intervals (Clopper-Pearson intervals) as each Semantic Exact Match is a binary outcome (0 or 1 for each example), and unlike bootstrapping this also supports the common case for our RASP program of 100% accuracy, which occurs in all but one split for our model, where resampling will not help us estimate uncertainty in bootstrapping, but the Beta distribution will still give us confidence bounds that depend on the sample size.

In developing our RASP program[59], when we find the right index of a verb relation (like agent, theme, or recipient), we found it was necessary to skip any noun phrases preceded by a preposition ("in", "on", "beside")[60].[61]

---

[57] Complement phrase support has been added in https://github.com/willy-b/learning-rasp/pull/7 , and the model scores 100% semantic exact match (95% confidence interval: 99.88% to 100.00%) on the full ReCOGS_pos Test set, and 100% exact string match (same confidence interval) on the dev set (test set results pending).

[58] https://github.com/IBM/cpg/blob/c3626b4e03bfc681be2c2a5b23da0b48abe6f570/src/model/cogs_data.py#L523

[59] https://github.com/willy-b/learning-rasp/blob/16a8e154b025e91c8e56965a1d475e49f69ebdbd/word-level-pos-tokens-recogs-style-decoder-loop.rasp#L934

[60] RASP code in Appendix 2: RASP for relation right index ignoring distractor "pp np"

[61] Otherwise, when modifying a simple sentence like "The cake burned" with a preposition to "The cake on the plate burned" we would switch the theme from the cake to the plate by accident. This cake example is the infamous obj pp to subj pp example, where training a Transformer successfully to represent the semantics of sentences like "John ate the cake on

Since in the RASP program both this and subject prepositional phrase modification require the same rule ignoring the "pp np" when finding right index candidates for agent, theme, recipient outputs, we hypothesized two things.

One, that 'np v_dat_p2 np pp np np'[62] generalization after training on 'np v_dat_p2 np np pp np' would be difficult like (Wu et al., 2024)'s obj_pp_to_subj_pp split.

Two, that augmenting the training data with v_dat_p2 recipient modified sentences like "Emma gave a friend in a house a cookie" might lead to crossover improved performance on the subject pp generalization (e.g. "The friend in a house smiled"; without adding any example of subjects with pp modification).

Thus we additionally train (Wu et al., 2024) baseline Transformers from scratch in two separate experiments to test these.

For one, 'np v_dat_p2 np pp np np'[63] generalization after training on 'np v_dat_p2 np np pp np' we train (Wu et al., 2024) Transformers with default configuration and default training data, then we add a new generalization split derived from (Wu et al., 2024)'s 'train.tsv' of 328 existing training examples where we have transferred the prepositional phrase from the theme to the recipient[64] in the 'v_dat_p2' sentence form with one prepositional phrase (see Appendix 4 for details and actual data sample).

For two, to see if augmenting the training data with v_dat_p2 recipient modified sentences has crossover benefit, we train separate default (Wu et al., 2024) Transformer but with their existing train.csv plus the additional theme-modified sentences mentioned above, same as those used for generalization testing in the other experiment; we confirm it does not know them, and separately on

fresh runs we try training on them to see if that can benefit other splits by teaching the Encoder-Decoder a general prepositional phrase handling rule (like ignore "pp np"). We then test on (Wu et al., 2024)'s normal test and generalization splits.

(Wu et al., 2024) baseline Encoder-Decoder Transformers trained from scratch are trained with random weight initialization multiple times with at least 10 different random seeds with all performance metrics averaged across runs with sample mean, sample size, and unbiased sample standard deviation reported. Statistical significance of comparisons between any Transformers performance sample means will be checked with Welch's unequal variance t-test with p-values greater than 0.05 definitely rejected, though stricter thresholds may be used where applicable. Confidence intervals will be reported using 1.96 standard errors of the sample mean as the 95% confidence interval for sample means with that N unless specified otherwise.

22 Acknowledgements

---

the plate" leads to a model that won't immediately generalize to being able to represent the meaning of "The cake on the plate burned" in logical form. In writing our RASP program this was observed as nothing to do with subjects or objects but just modifying noun phrases to the left of the part of speech (say a verb) they have a relationship with, instead of on the right side. For example, this also occurs in v_dat_p2 sentences like "Emma gave a friend a cookie" (agent, recipient, theme nps). It is obvious that modification of the theme with prepositional phrases is not going to disrupt parsing the sentence: "Emma gave a friend a cookie (modification modification ...)", whereas modifying the recipient, on the left, due to the asymmetry of prepositional phrases adding to the right, disrupts the sentence, rendering it unreadable in the limit of too many pps:

"Emma gave a friend (modification modification ...) a cookie" , in the limit of more modification, "a friend" cannot be associated with "a cookie".

[62]Being precise we only do 'np v_dat_p2 np_det pp np np' as per the grammar 'np_prop' cannot precede a prepositional phrase

[63]Restricted to 'np v_dat_p2 np_det pp np np' as per the grammar 'np_prop' cannot precede a prepositional phrase

[64]When the recipient is np_det, not np_prop; and we confirm it is within the grammar by reparsing with the Lark parser on the original grammar rules.